



---

# American Express-Default Prediction

Predict if a customer will default in the future



**team\_ amazon**

정주희

**Team Leader**

- XGBoost, LightGBM Modeling
- PPT
- Announcement

손희경

**Team Member\_1**

- Random Forest Modeling
- EDA
- Organize Data

최유림

**Team Member\_2**

- CatBoost Modeling
- EDA
- Organize Data

## >> Table of contents

- 1 Competition Description
- 2 Data Preprocessing and EDA
- 3 Insight about Machine Learning Model
- 4 Conclusion

# 1



## Competition Description

# Discount Revenue(60%)

: 카드 결제 수수료

American Express Business Model

## American Express Closed Loop

Card network, card issuer, merchant acquirer & processor



revenues

Discount Revenue  
Net Card Fees  
Other Fees  
Net Interest Income



## Discount Revenue(60%)

카드 결제 수수료

“

American Express Business Model

Whether out at a restaurant or buying tickets to a concert, modern life counts on the convenience of a credit card to make daily purchases. It saves us from carrying large amounts of cash and also can advance a full purchase that can be paid over time. How do card issuers know we'll pay back what we charge? That's a complex problem with many existing solutions—and even more potential improvements, to be explored in this competition.

”



Discount Revenue(60%)

카드 결제 수수료

“

American Express Business Model

How do card issuers know  
we'll pay back what we charge?



CARD ISSUERS



MERCHANTS

REVENUES



”





## anonymous

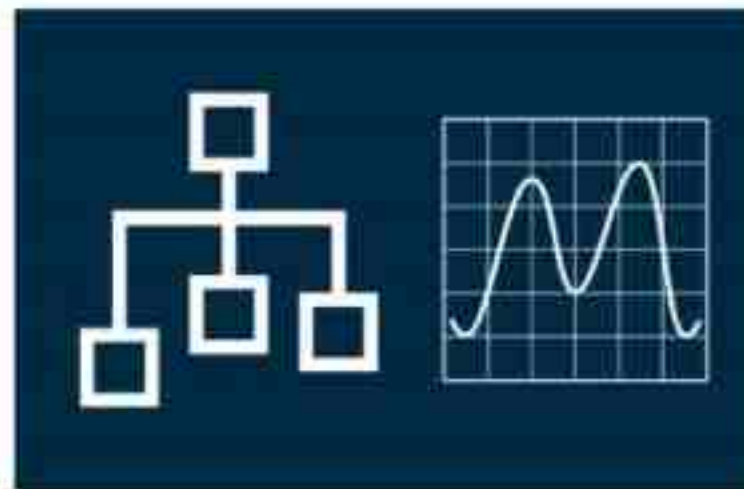
anonymized customer  
profile information



## time-series data

The target binary variable is calculated  
by observing 18 months performance window  
after the latest credit card statement

target=1 : does not pay due amount in 120 days  
after their latest statement date

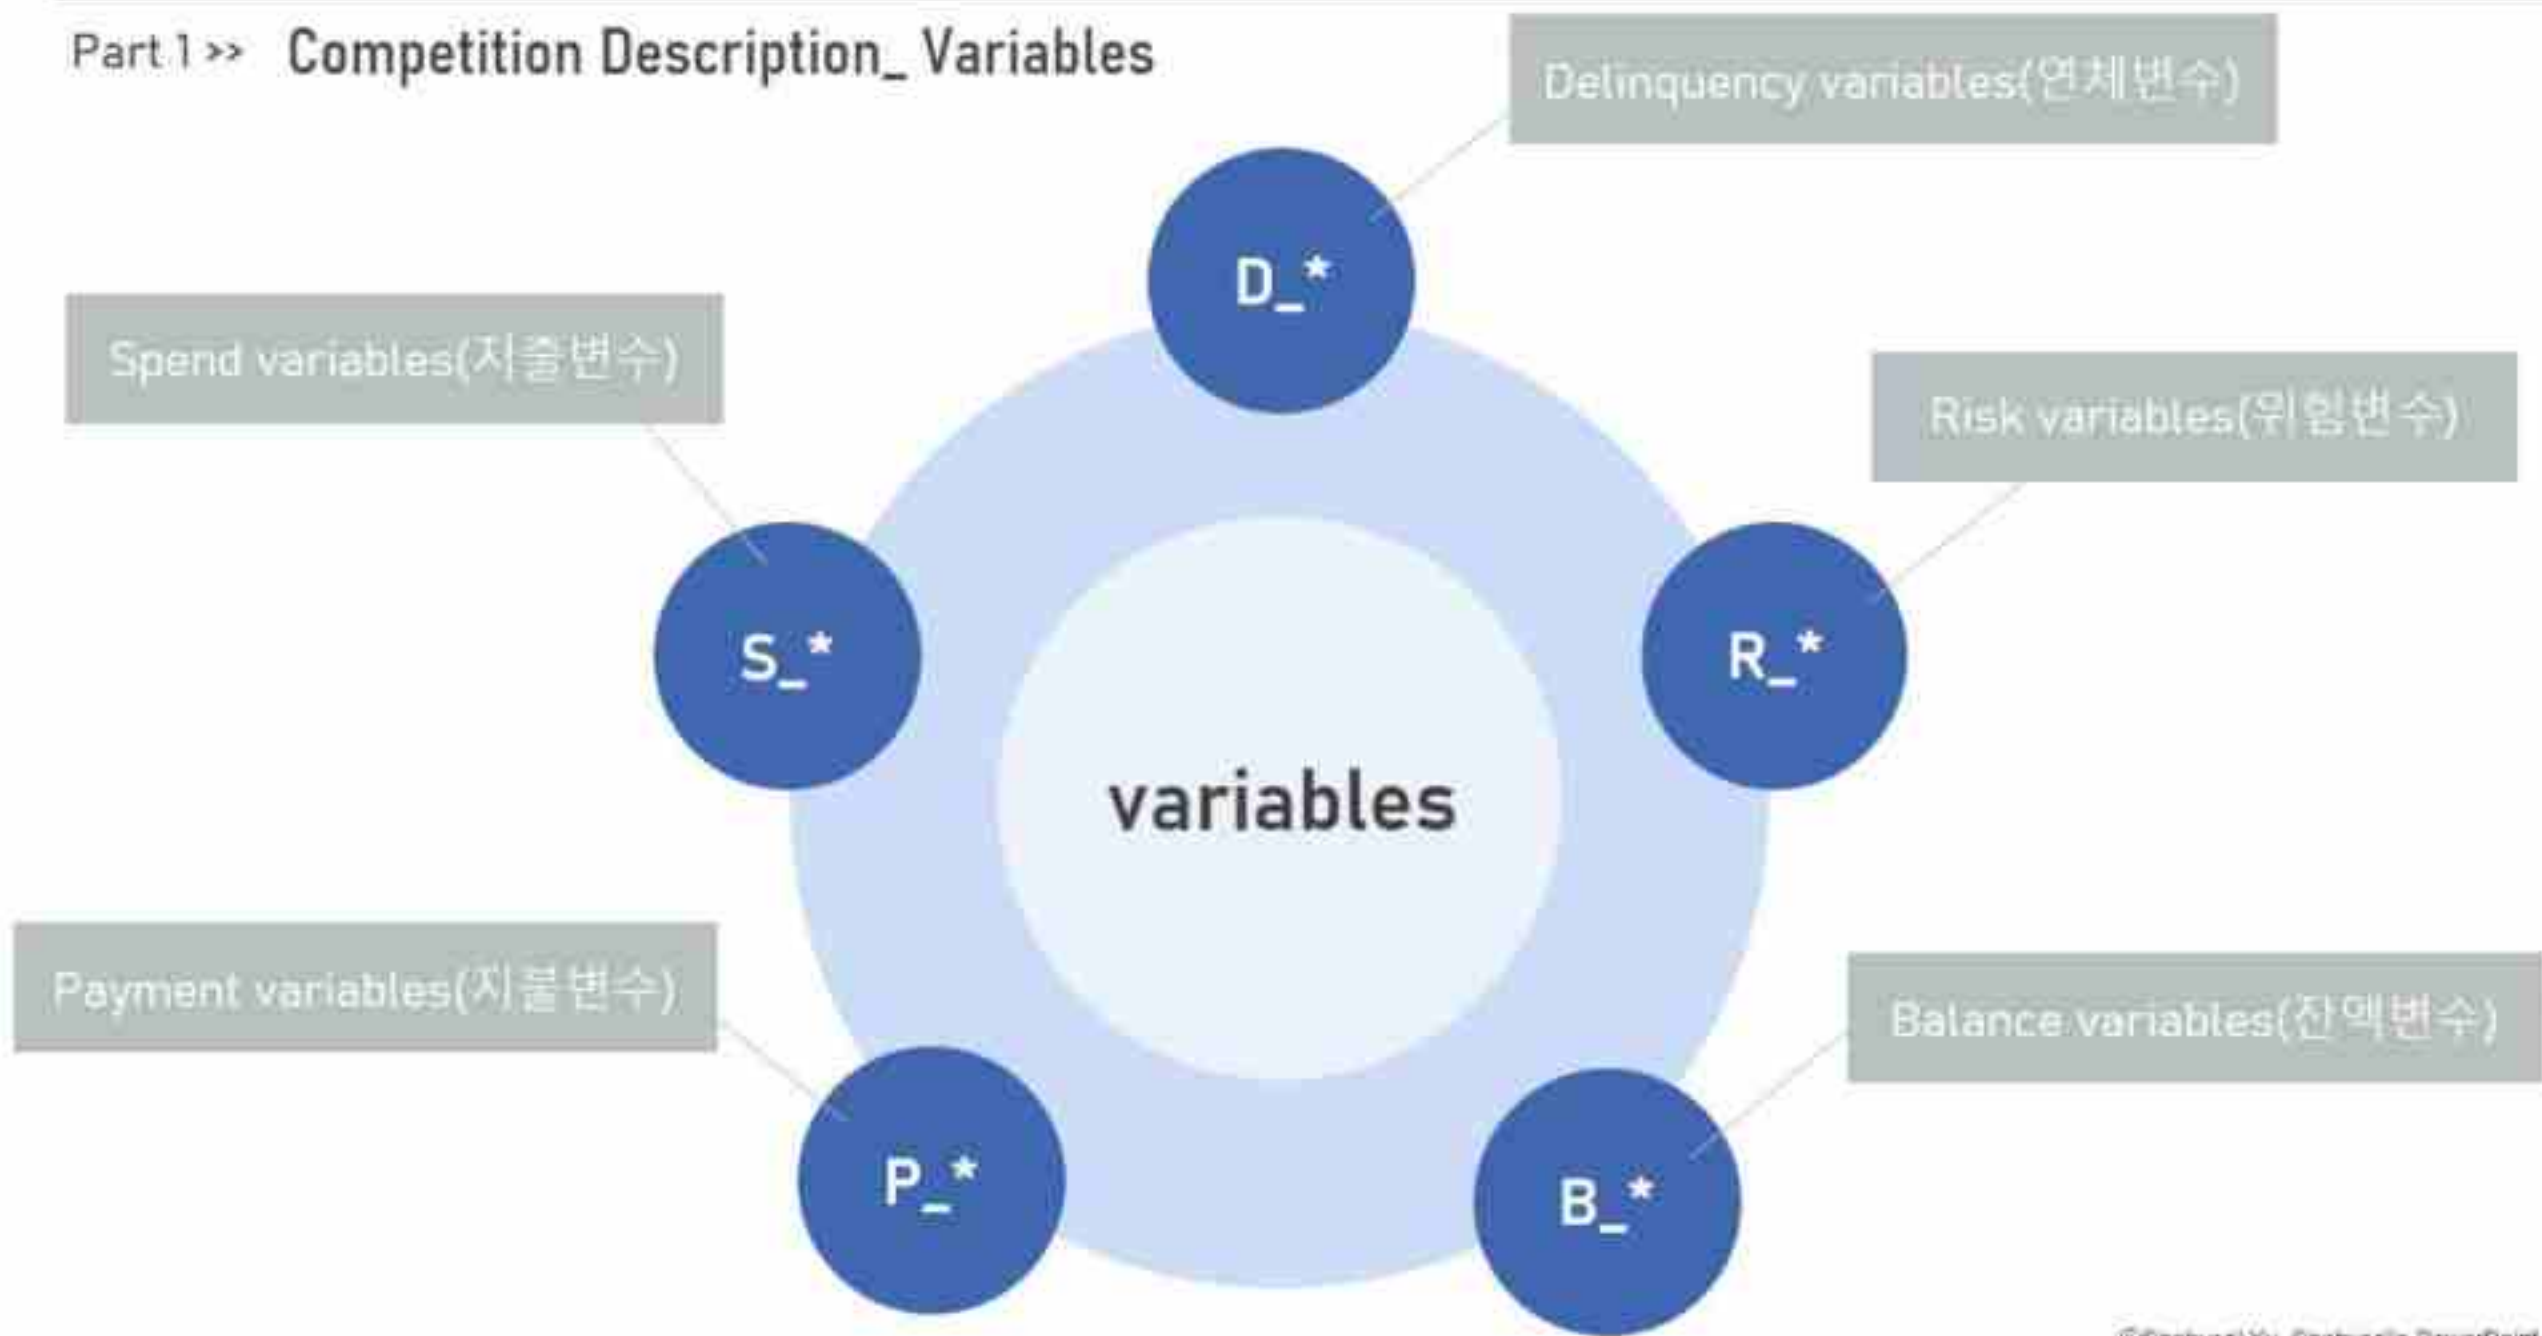


## categorical/numerical

C : 'B\_30', 'B\_38', 'D\_114', 'D\_116', 'D\_117', 'D\_120',  
'D\_126', 'D\_63', 'D\_64', 'D\_66', 'D\_68'

predict 'target = 1' for each customer\_ID

Part 1 >> Competition Description\_Variables



## Understand Customer Data\_ Unique Customer

In [7]:

```
unique_customer_count = len(train.groupby("customer_ID")["customer_ID"].count())  
print("unique customer data in training data -->", unique_customer_count)  
unique_customer_count_test = len(test.groupby("customer_ID")["customer_ID"].count())  
print("unique customer data in test data -->", unique_customer_count_test)
```

```
unique customer data in training data --> 458913  
unique customer data in test data --> 924621
```

# 2



## Data Preprocessing and EDA

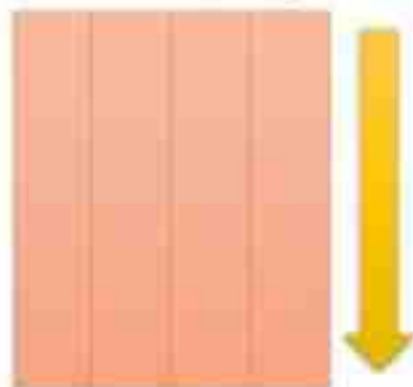
## Part 2 >> Exploratory Data Analysis



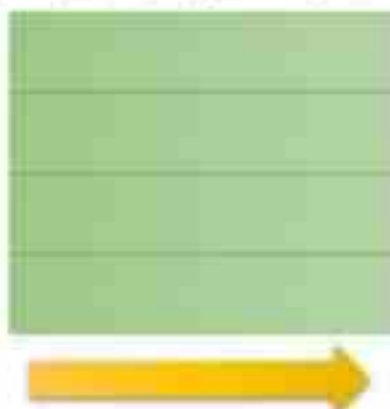
## Part 2 >> Exploratory Data Analysis\_ How to reduce Data size?



열 기반 압축(Parquet)

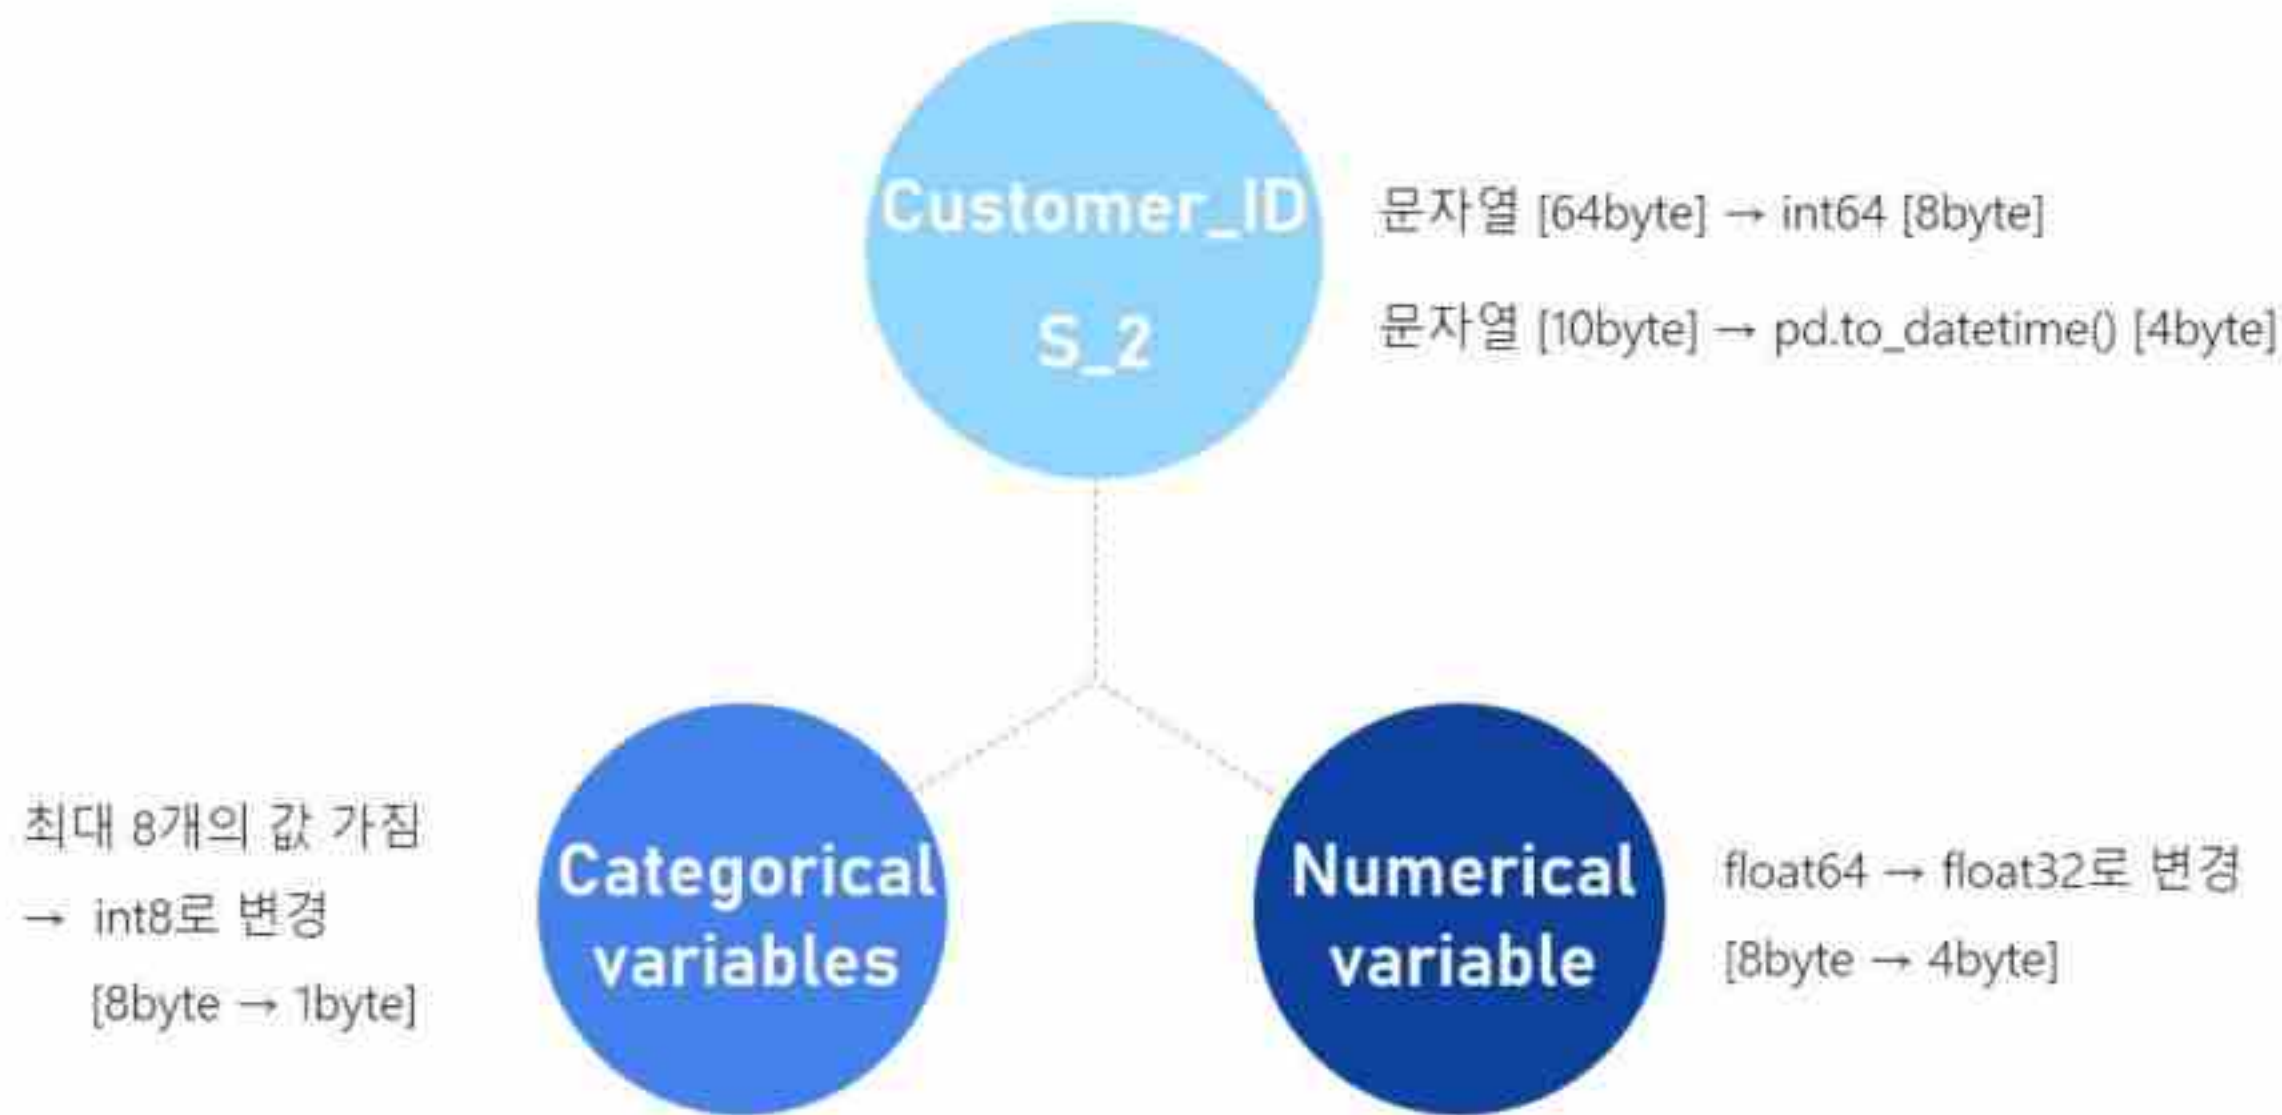


행 기반 압축(전통적 방식)



- ✓ 행 기반 압축 → 데이터 압축률 우수
- ✓ 특정 column의 데이터만 읽고 처리 가능  
→ 데이터 처리에 들어가는 자원 절약됨
- ✓ column에 동일한 dtype이 저장  
→ column별로 적합한(데이터형에 유리한) 인코딩 사용가능
- ✓ 특정 column을 선택해서 가져오는 형식  
→ 선택하지 않은 column 데이터는 I/O가 발생하지 않음  
→ 시간, 메모리 사용량 줄일 수 있음
- ✓ 파일 형식이 dtype을 기억  
→ int64에서 int8로 downcast하면 다음에도 파일을 int8로 불러옴





## Part 2 >> Exploratory Data Analysis\_ Missing Value

### Verification\_Missing Value



“

Verification\_ Missing Value

Amount of Missing Data

결측치가 매우 많음. 처리 필요

→ 결측치 비율 80%를 넘는 feature: 23개, 이들 제거함

”

“

### Verification\_Missing Value

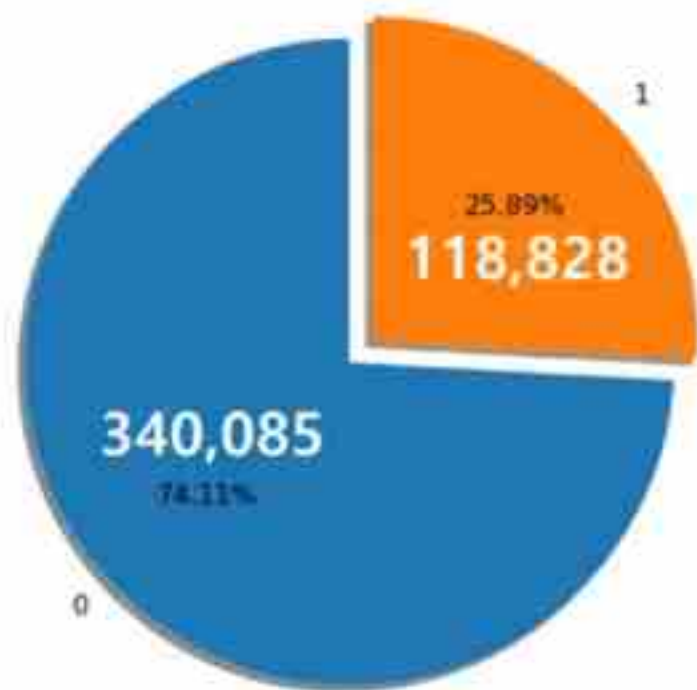
```
# Removing Columns With NaNs Rate Higher Than Threshold
nan_pct_threshold = 80
to_remove_cols = list(nan_values_pct[nan_values_pct > nan_pct_threshold].index)
print(f"Columns With NaN Values Rate > {nan_pct_threshold}%: {len(to_remove_cols)} Columns")
train_data = train_data.drop(columns=to_remove_cols).reset_index(drop=True)
```

Columns With NaN Values Rate > 80%: 23 Columns

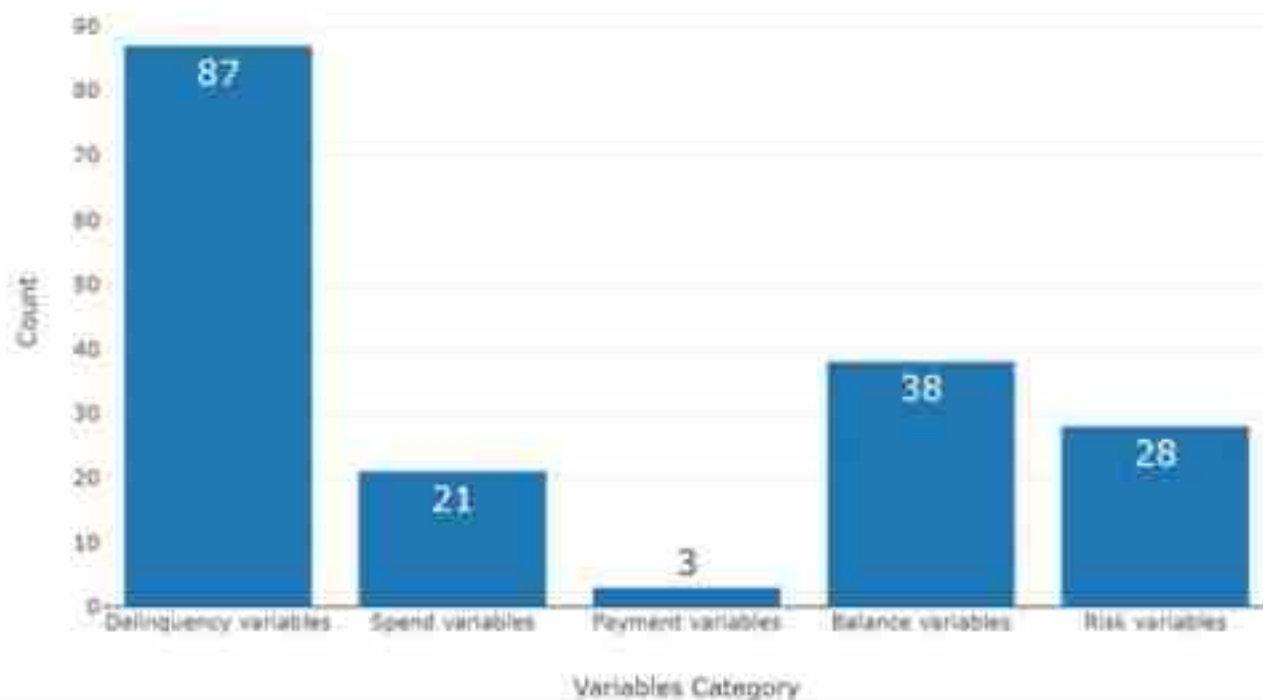
”

### Verification\_After Missing Value Processing

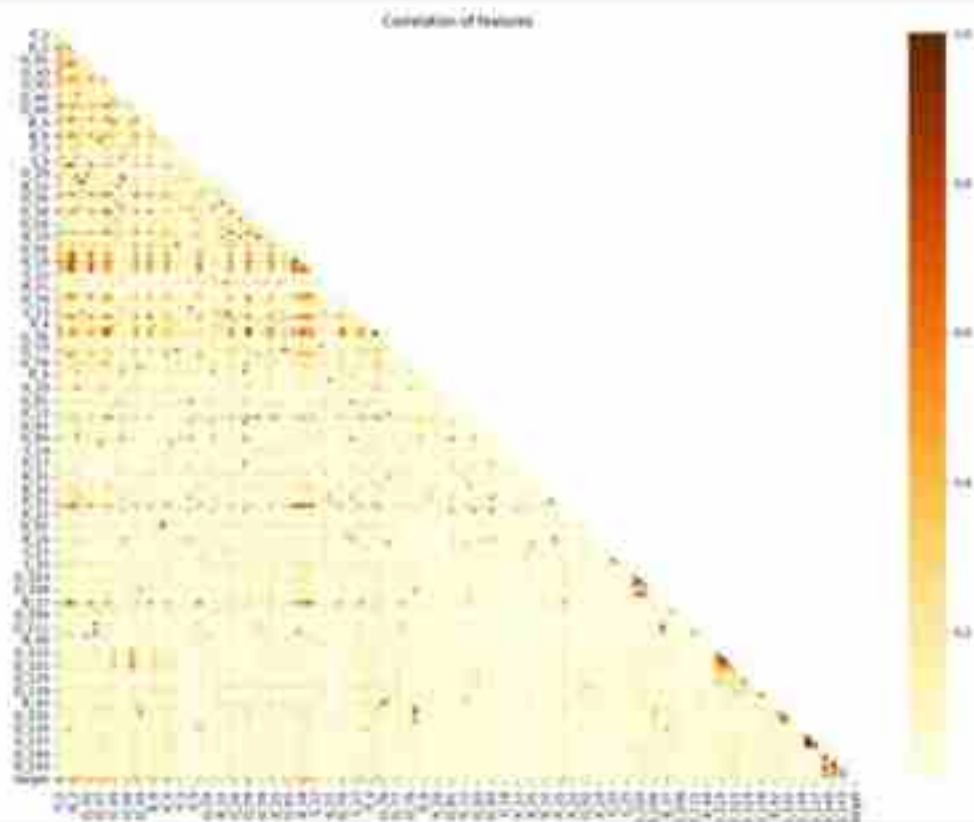
Target Variable Distribution



Variables Count By Category



Verification\_ Correlation between variables(heatmap)



insight

Feature 상관관계 파악필요  
→ Heatmap이용해 시각화

상관관계가 뚜렷하지 않음

일부 색이 진한 점들은 다시  
결과값 unstack해야 함

## Part 2 >> Exploratory Data Analysis\_ Statistical Values

### Verification\_ Correlation between Unstacked Variables

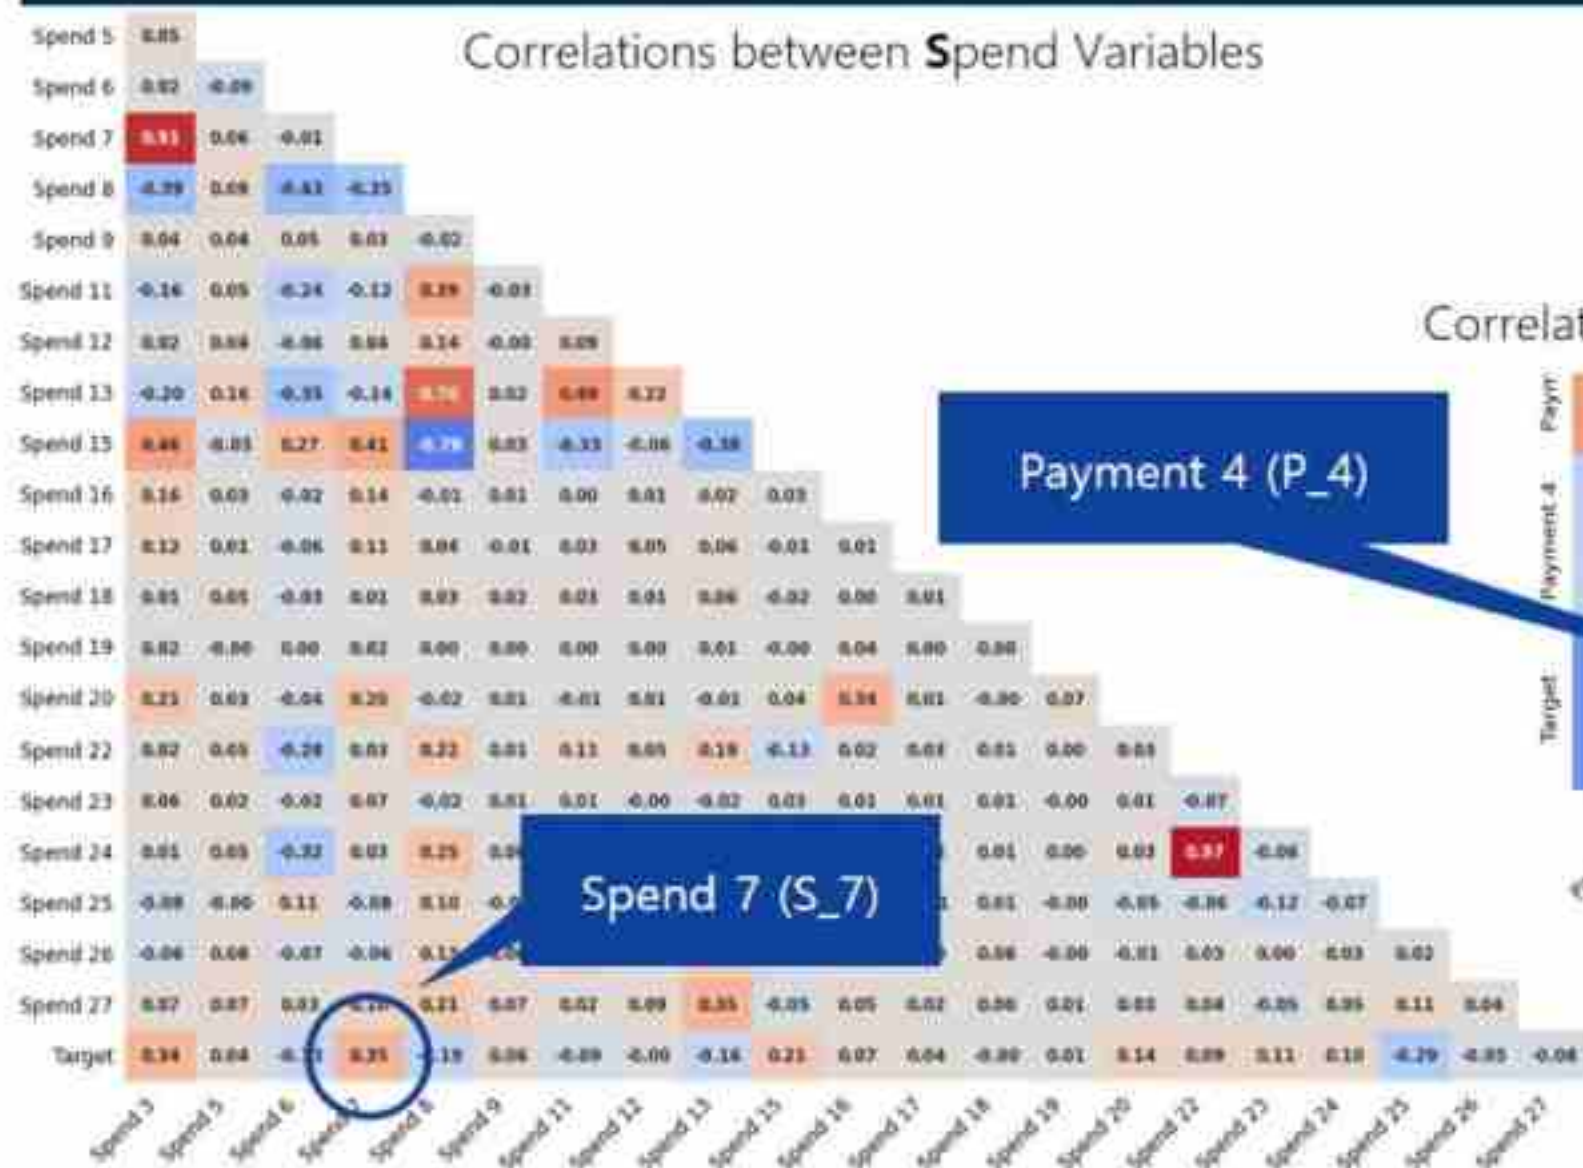
```
In [38]: unstacked = correlations.unstack()
unstacked = unstacked.sort_values(ascending=False, kind="quicksort").drop_duplicates().head(25)
unstacked
```

V_4	P_3	
0.143	0.144	0.999775
0.77	0.82	0.998734
0.138	0.140	0.998629
	0.143	0.998261
0.145	0.141	0.997764
0.11	0.1	0.996470
0.22	0.7	0.995641
0.111	0.118	0.994961
0.1	0.37	0.993289
0.37	0.11	0.991103
0.19	0.74	0.987034
0.127	0.120	0.983981
0.22	0.24	0.978798
0.38	0.70	0.925743
	0.74	0.920571
0.138	0.127	0.918382
0.14	0.15	0.913476
0.20	0.2	0.912721
0.126	0.128	0.908674
0.127	0.126	0.905224
0.7	0.3	0.903646
0.125	0.138	0.903043
0.40	0.78	0.898608
0.28	0.16	0.898182

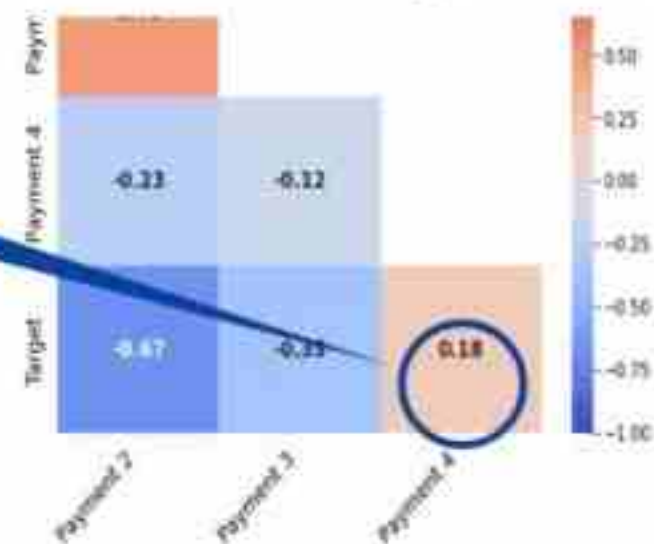
dtypes: float64 3



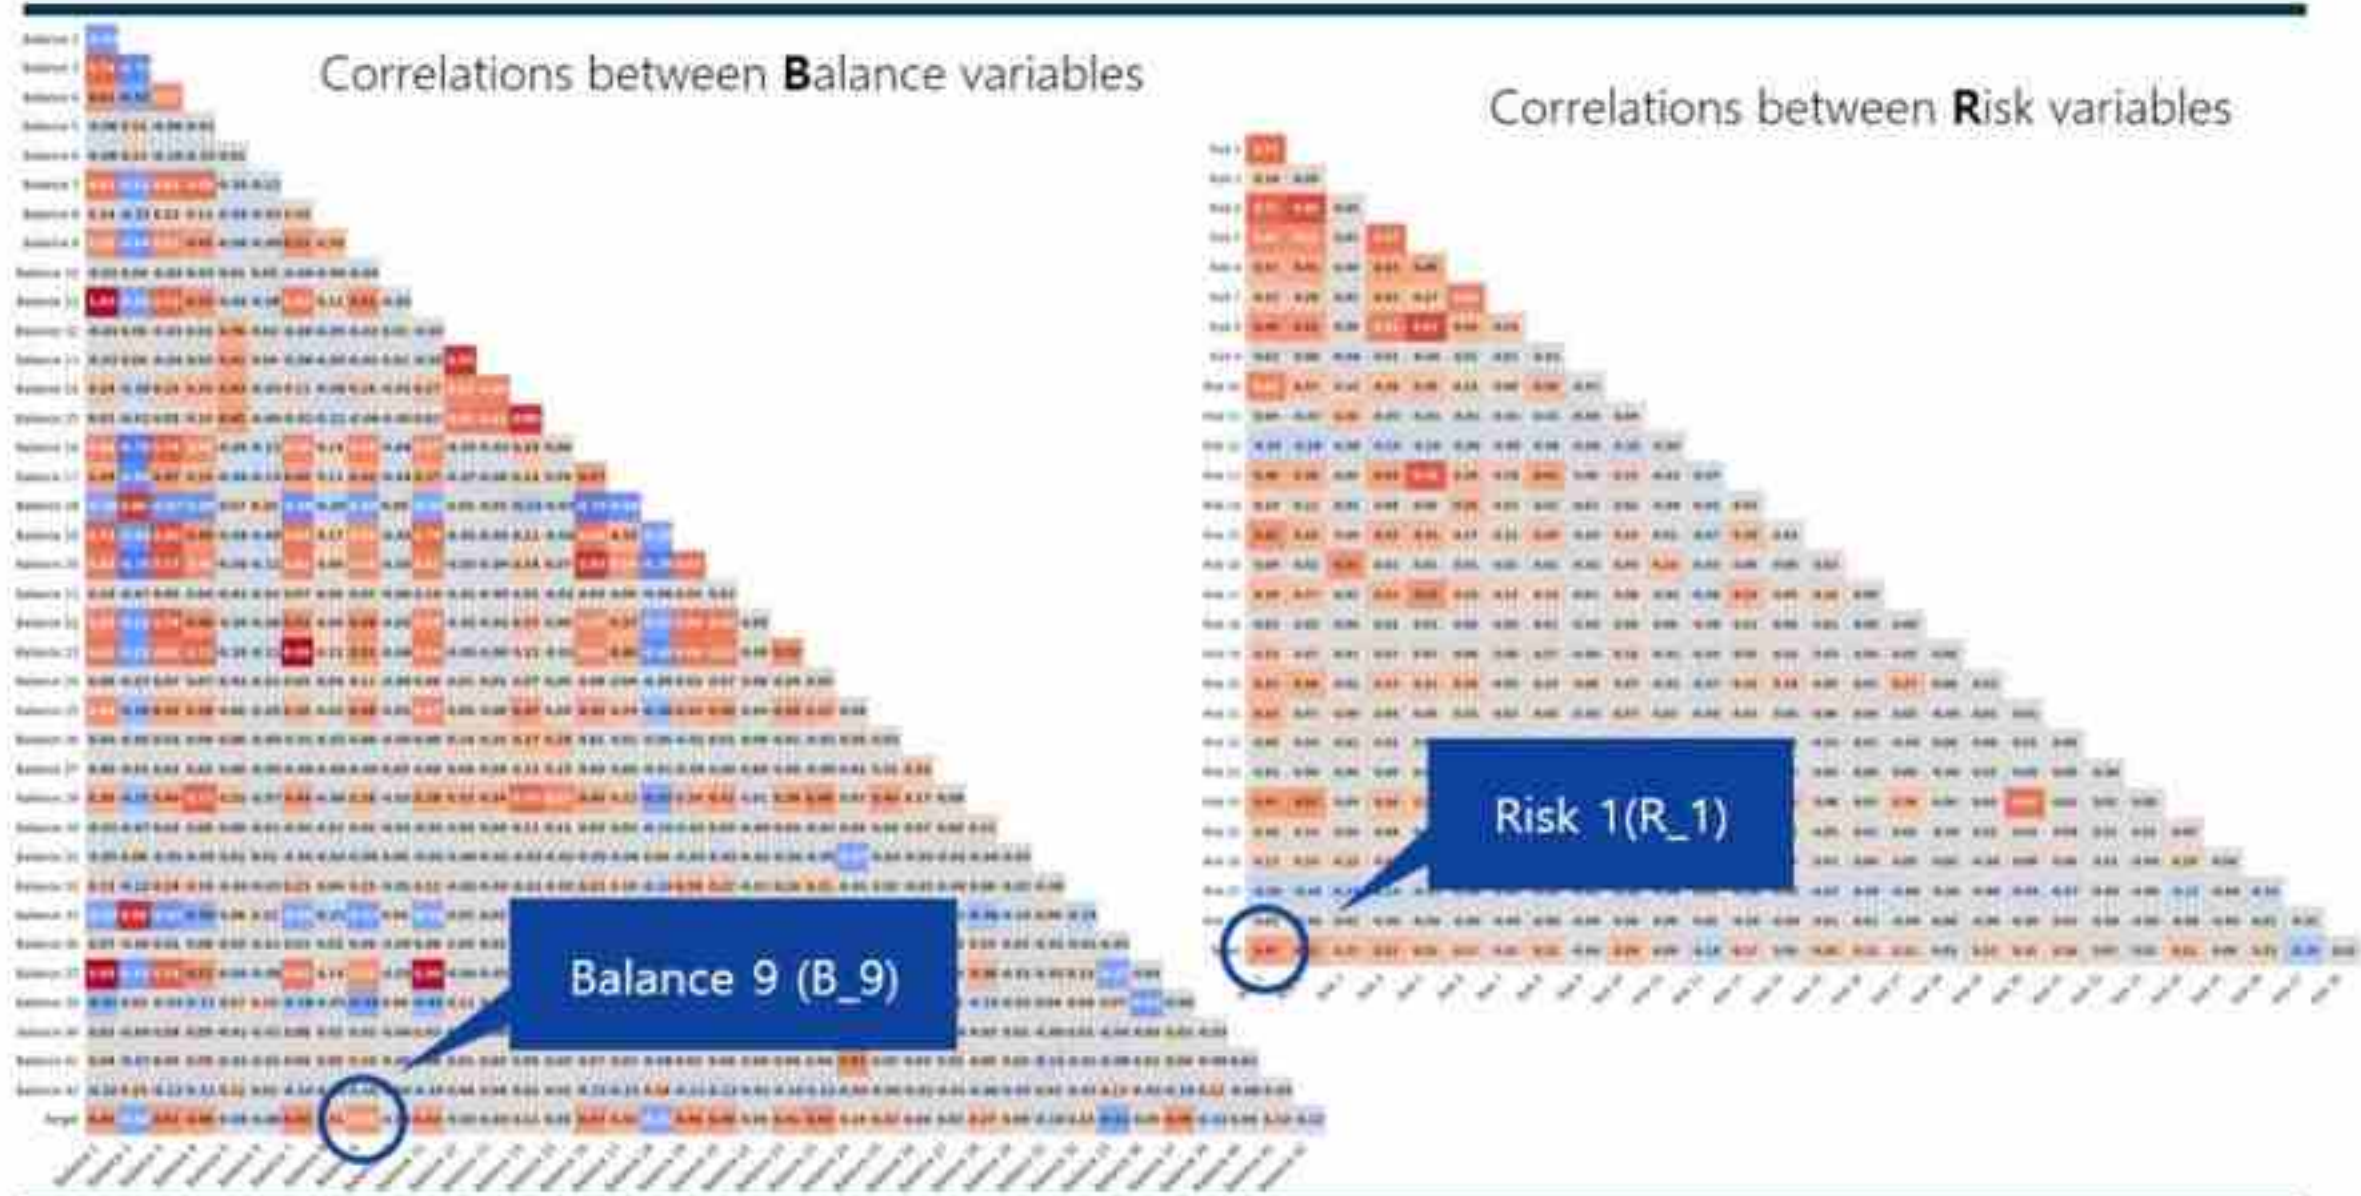
## Part 2 >> Exploratory Data Analysis\_ Statistical Values



Correlations between **Payment** variables



## Part 2 >> Exploratory Data Analysis\_ Statistical Values



## Part 2 >> Exploratory Data Analysis\_ Statistical Values

Verification\_ Statistical value (count,mean,std,min,max,Quantiles)\_target '0'

```
ex_customer_data[ex_customer_data["target"] == 0][b_cols[:10]].describe()
```

	B_1	B_2	B_3	B_4	B_5	B_6	B_7	B_8
count	9032.000000	9032.000000	9032.000000	9032.000000	9032.000000	9032.000000	9032.000000	8.989000e+03
mean	0.081276	0.719796	0.081871	0.130928	0.097953	0.161655	0.133657	3.310980e-01
std	0.156970	0.355266	0.181874	0.182206	0.331847	0.319970	0.192023	4.688353e-01
min	-0.141469	0.000164	0.000003	0.000017	0.000007	-0.000552	-0.096913	2.764867e-07
25%	0.007554	0.620152	0.004455	0.019144	0.007483	0.037204	0.024911	3.803379e-03
50%	0.021453	0.817121	0.008435	0.055663	0.017671	0.140455	0.045870	7.330273e-03
75%	0.061725	1.003689	0.041836	0.161069	0.072235	0.199933	0.156390	1.002241e+00
max	1.320823	1.009999	1.171260	1.283849	12.974426	20.331217	1.252293	1.010181e+00



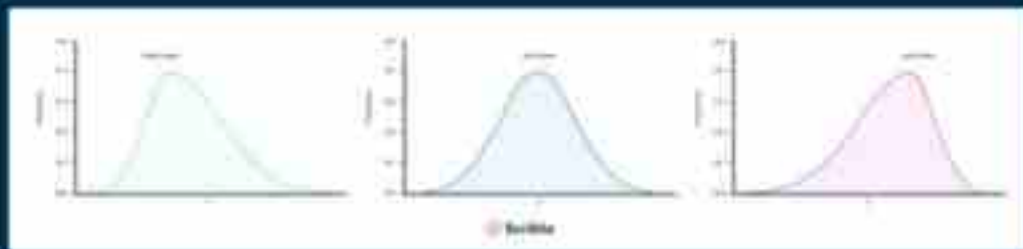
Verification\_ Statistical value (count,mean,std,min,max,Quantiles)\_target '1'

```
ex_customer_data[ex_customer_data["target"] == 1][b_cols[:10]].describe()
```

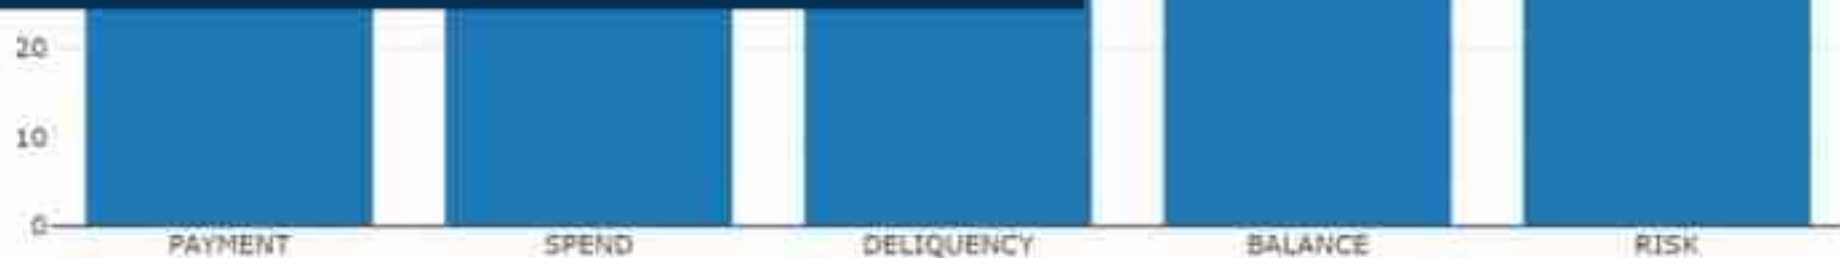
	B_1	B_2	B_3	B_4	B_5	B_6	B_7	B_8
count	3030.000000	3030.000000	3030.000000	3030.000000	3030.000000	3030.000000	3030.000000	3022.000000
mean	0.257518	0.299894	0.298945	0.328520	0.031458	0.043737	0.355820	0.734537
std	0.279484	0.375838	0.300259	0.296676	0.084002	0.085808	0.252973	0.444182
min	-0.046796	0.000034	0.000085	0.000082	0.000006	-0.002122	0.000510	0.000006
25%	0.054462	0.028687	0.036761	0.120258	0.007023	0.009662	0.159904	0.009335
50%	0.136546	0.066485	0.220646	0.251282	0.011993	0.017531	0.309124	1.003016
75%	0.378923	0.610611	0.455126	0.450153	0.023515	0.038585	0.510301	1.006530
max	1.323411	1.009960	1.258546	2.187350	2.188328	1.926984	1.252394	1.010065

## Part 2 >> Exploratory Data Analysis\_ Outlier detection for skewed data

Verification\_Skewness Rate for each Variable Category

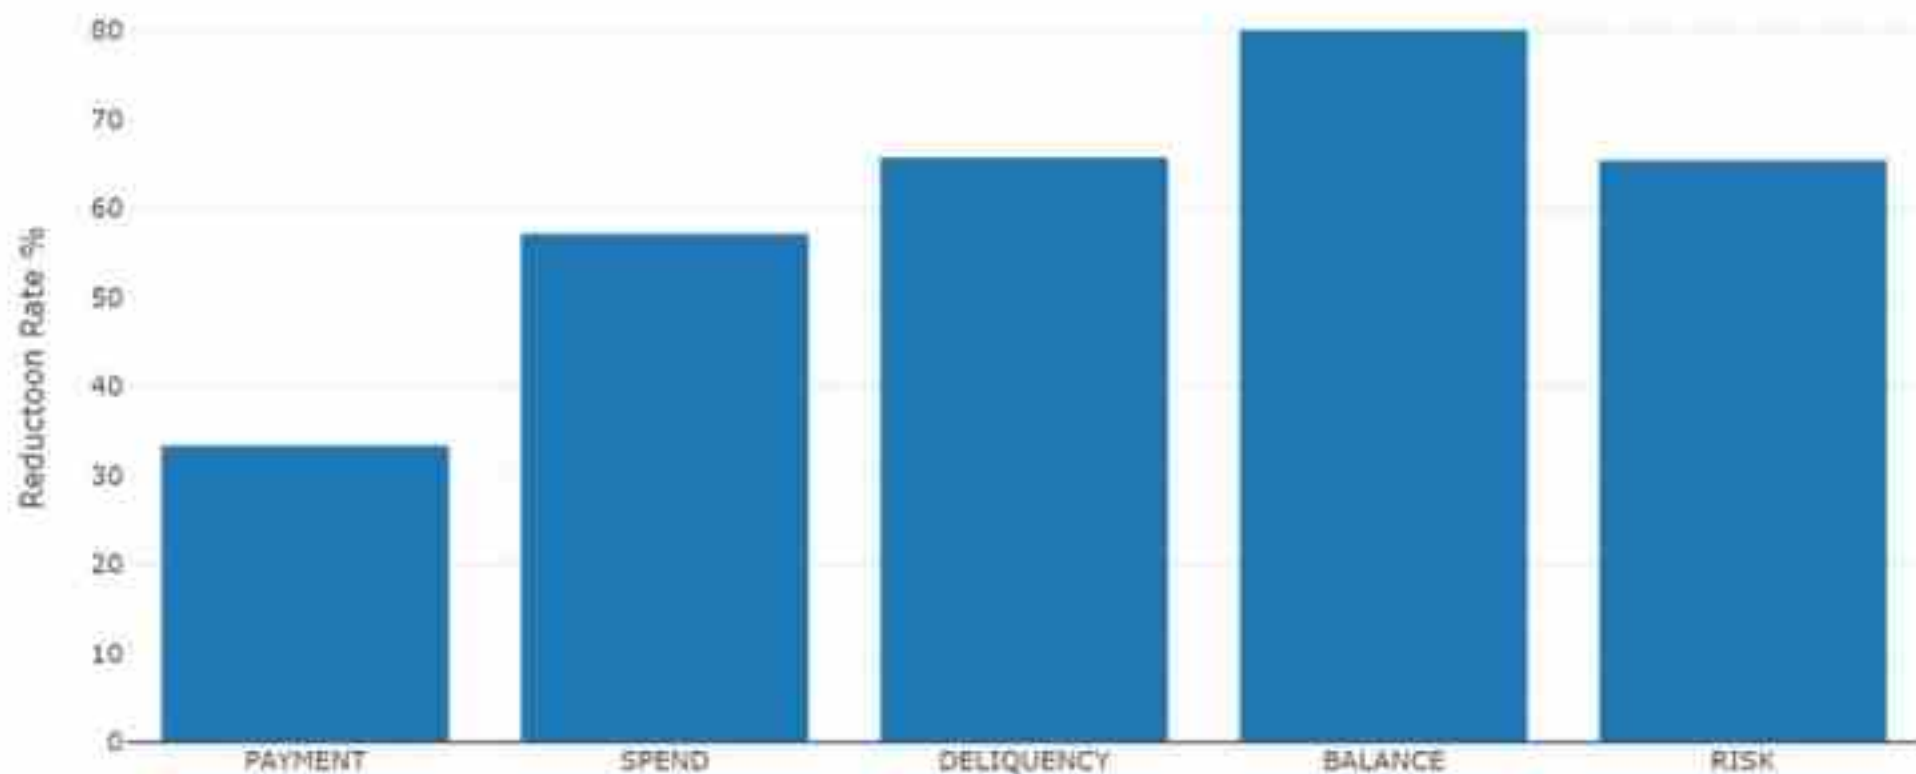


- deviate from the symmetrical bell curve
- Outliers(nonsense) → skewness



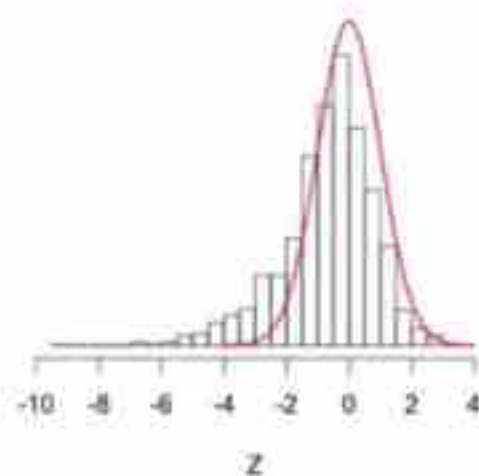
## Part 2 >> Exploratory Data Analysis\_ Outlier detection for skewed data

Verification\_ Skewness Reduction Rate for each Variable Category

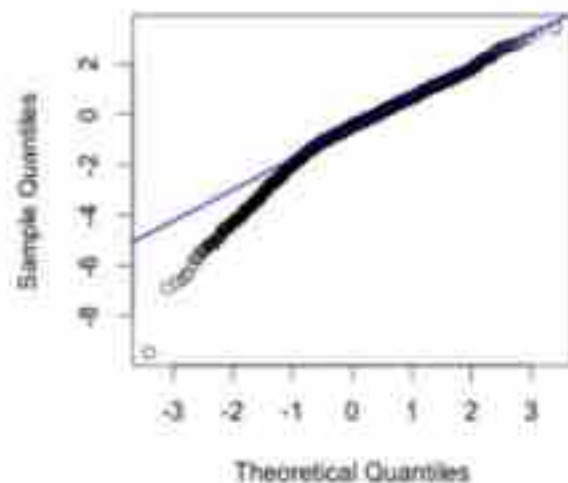


Verification\_Quantile-Quantile plot

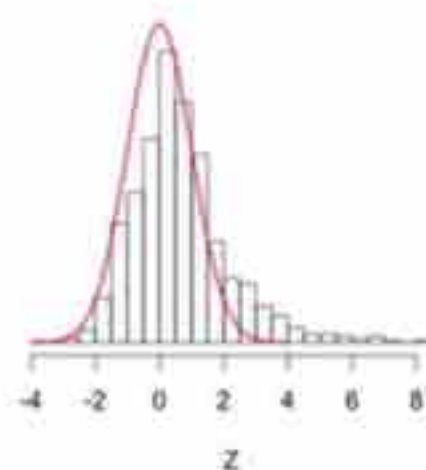
Skewed Left



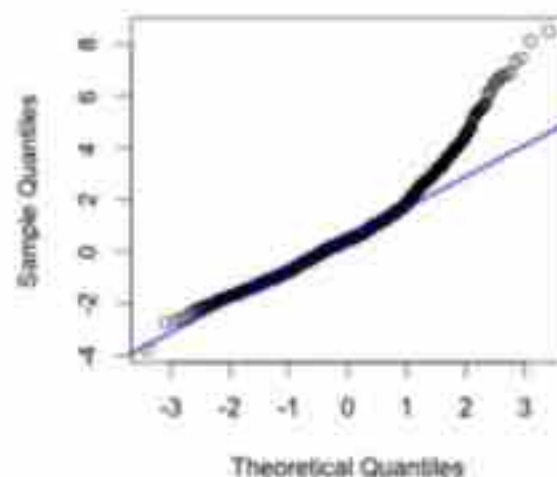
Normal Q-Q Plot



Skewed Right

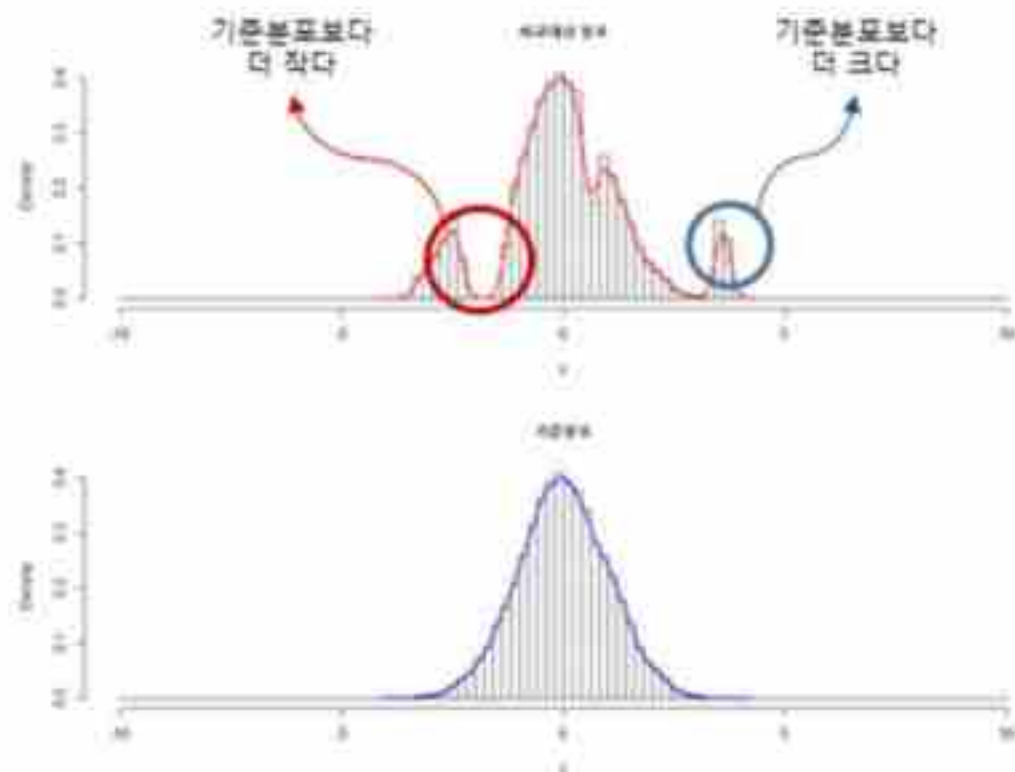
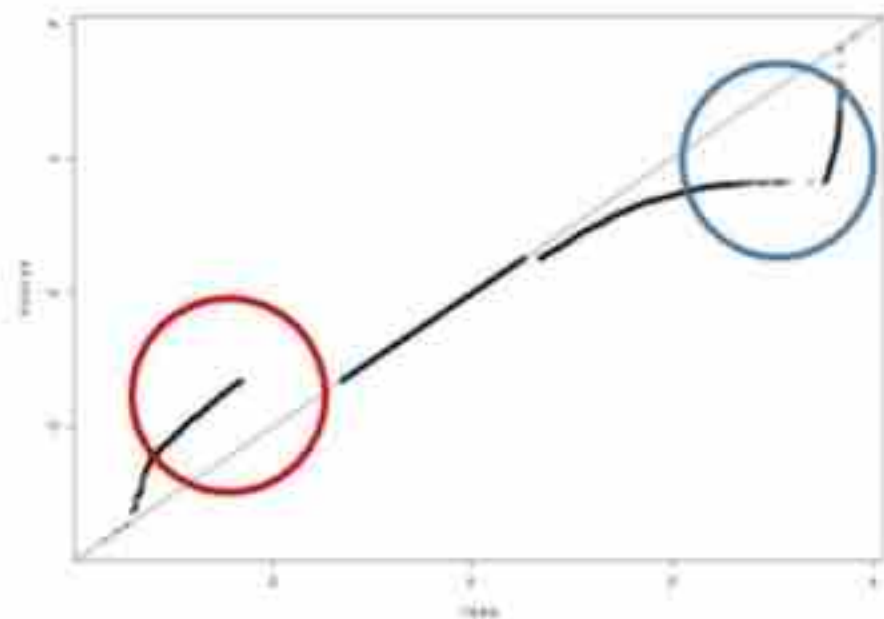


Normal Q-Q Plot



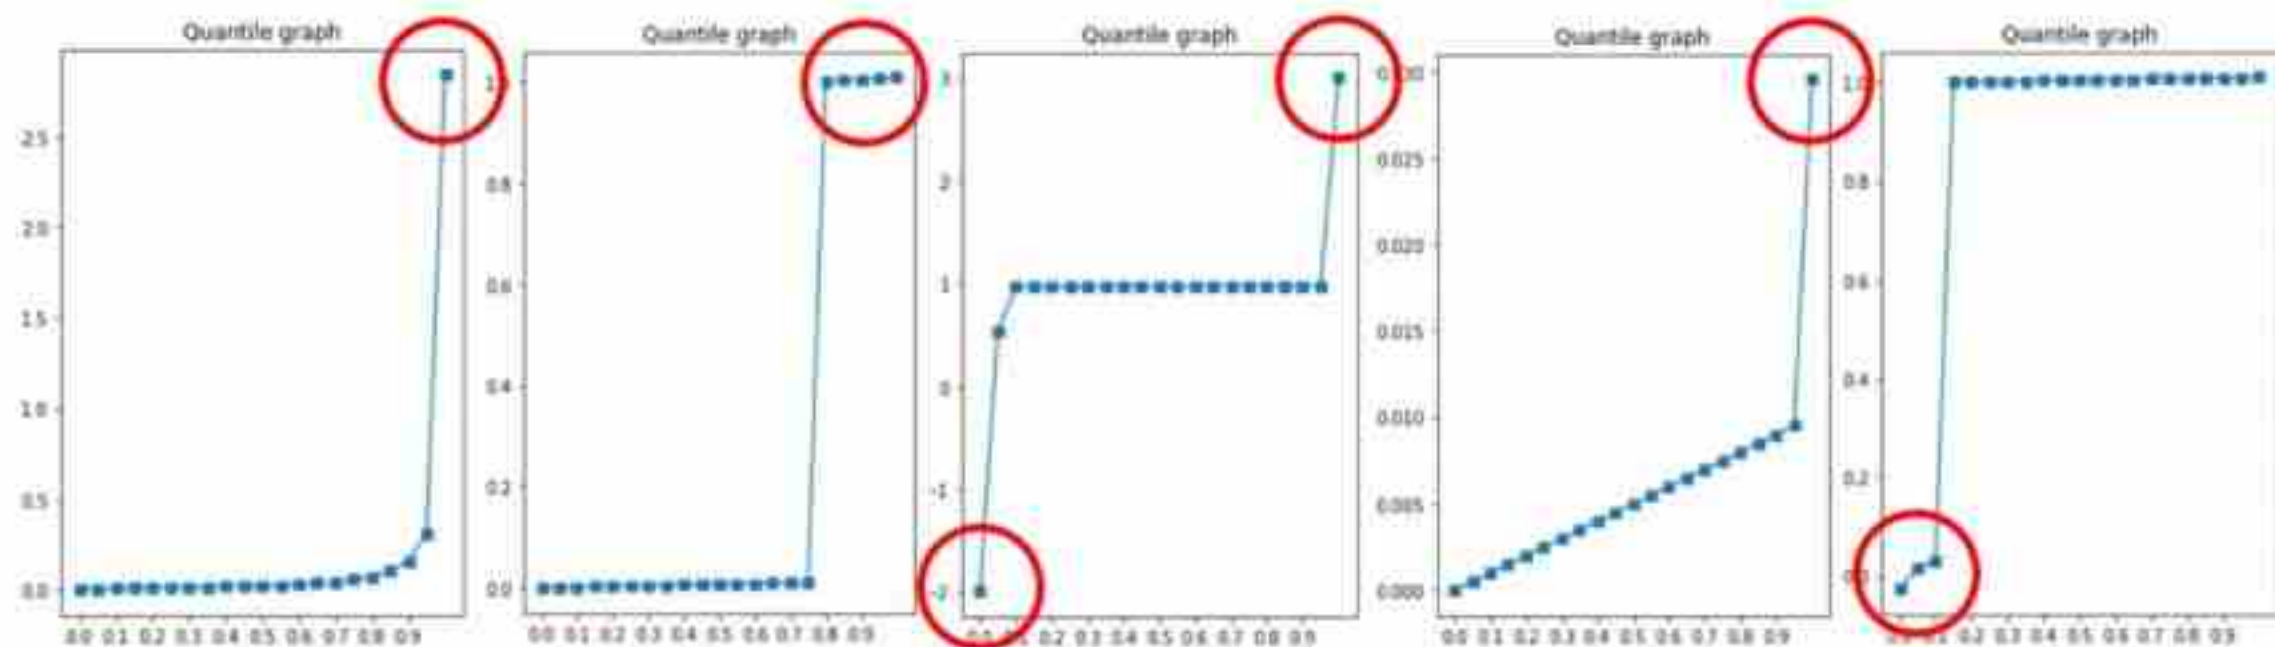


Verification\_Quantile-Quantile plot



## Part 2 >> Exploratory Data Analysis\_ Understanding Features with Quantiles

### Verification\_ Quantile graph for Outlier detection



Verification\_ Quantile graph for Outlier detection

“

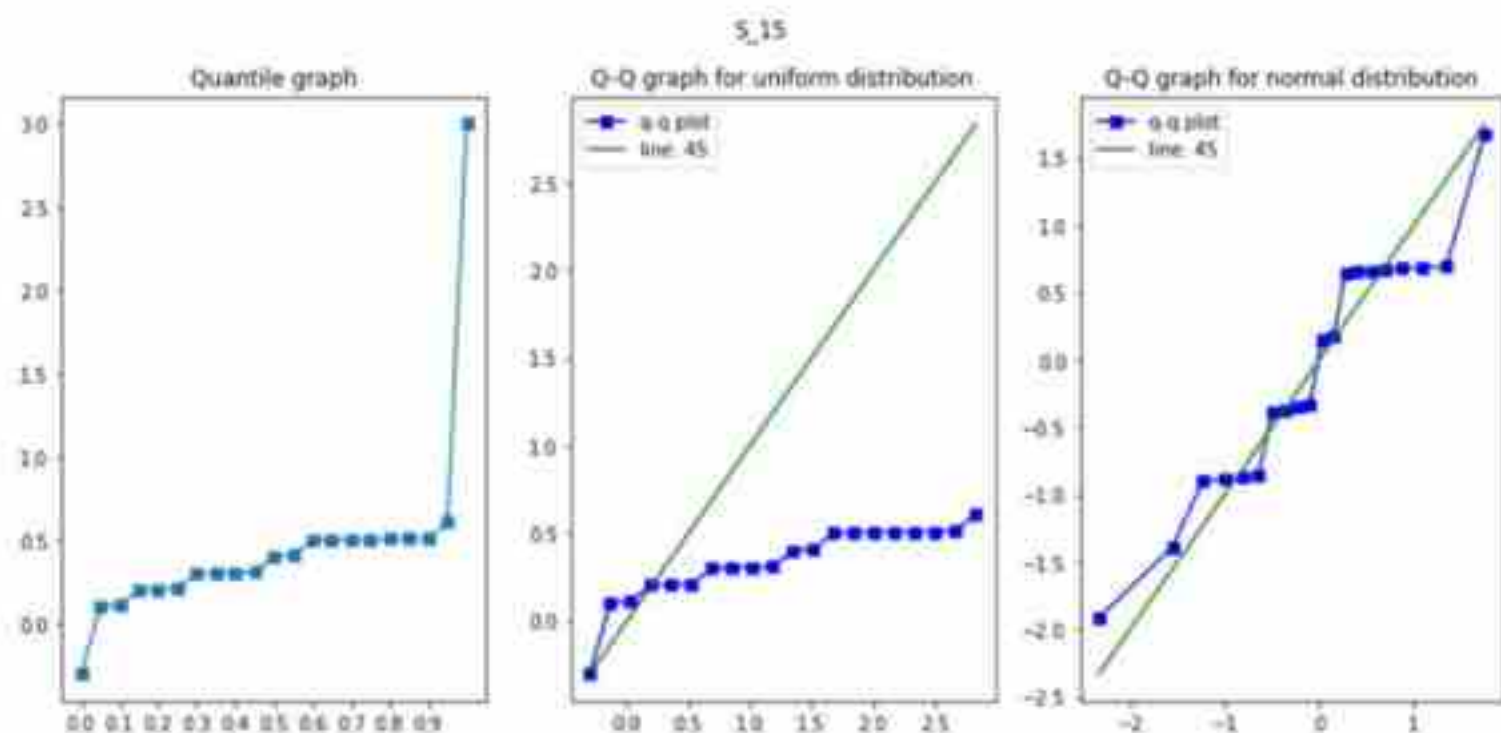


”

## Part 2 >> Exploratory Data Analysis\_ Understanding Features with Quantiles

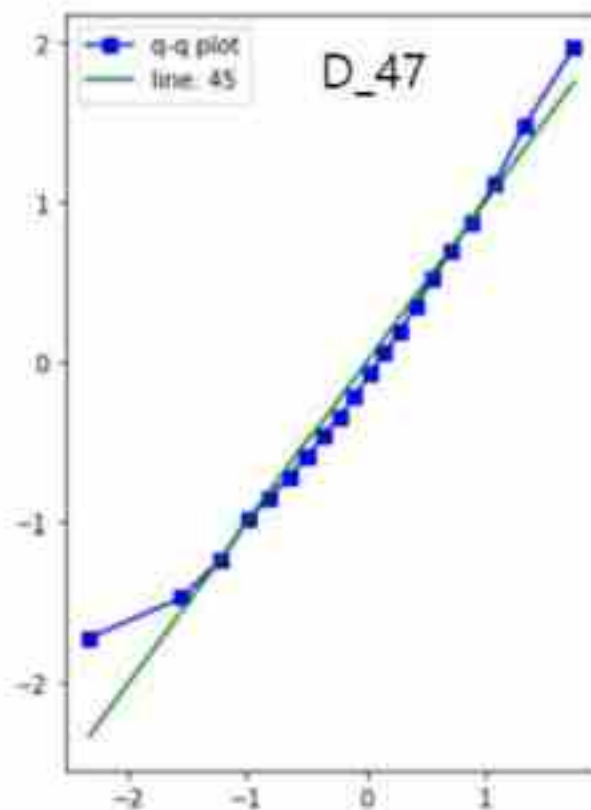
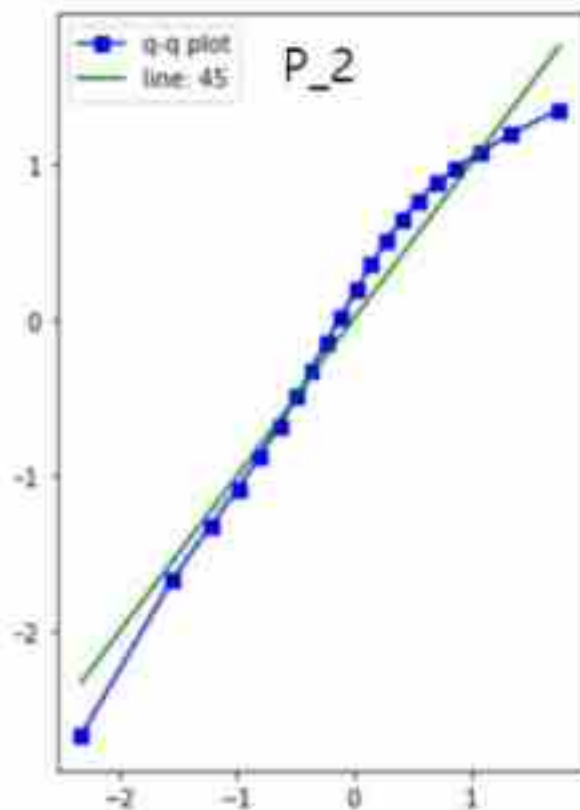
Verification\_ Highly discrete values

insight



S15는 최소 7개의 이산 값들  
가지는 것으로 보임

Verification\_ Q-Q graph for Normal Distribution



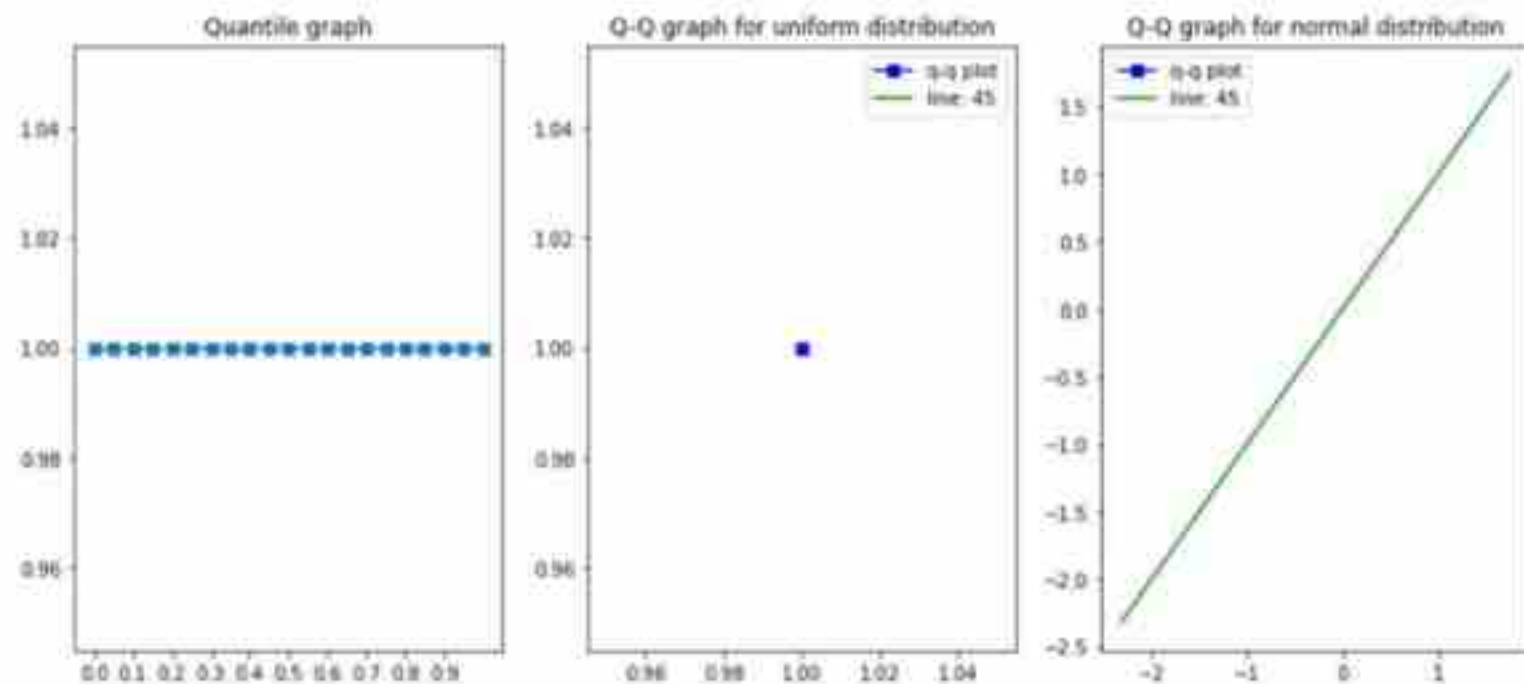
insight

P\_2, D\_47은  
가우시안분포(=정규분포)  
따르는 것으로 판단됨

Verification\_ Q-Q graph for some Features

insight

B\_31

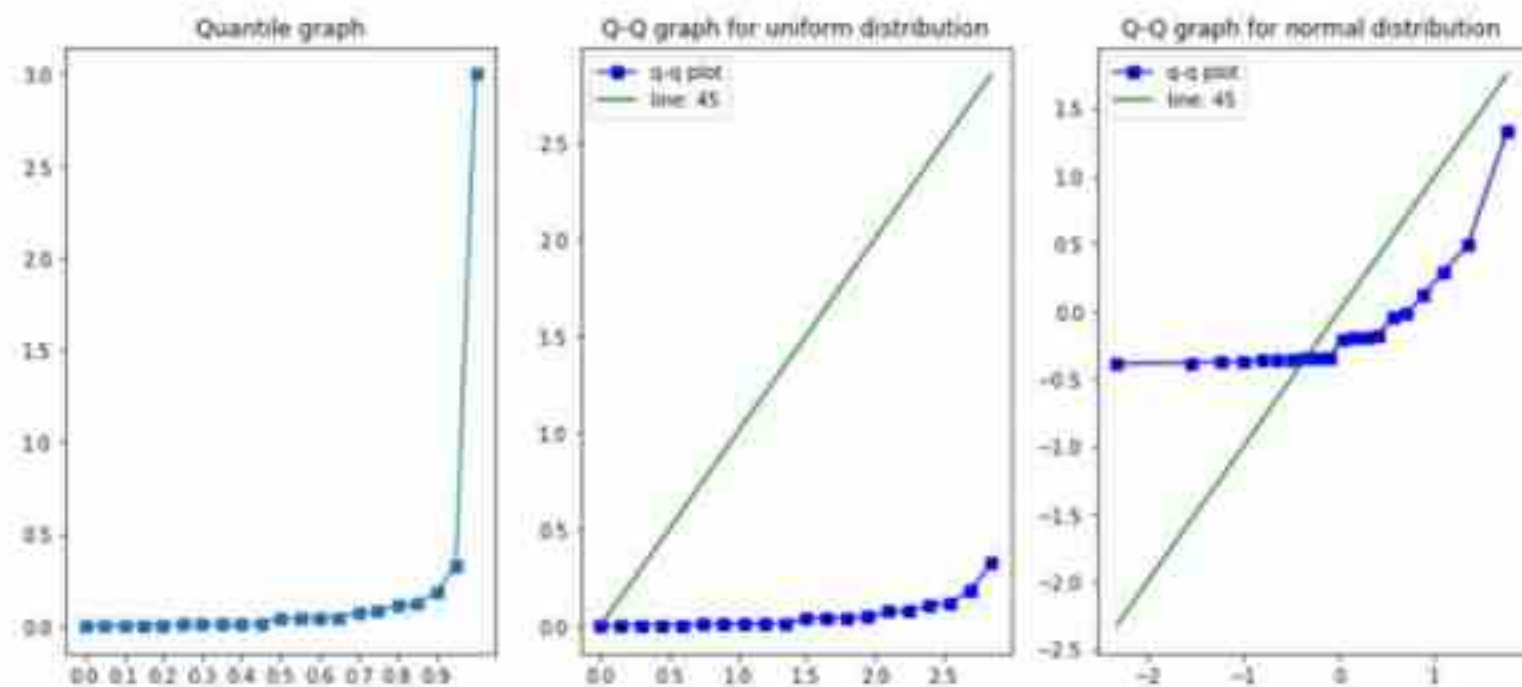


B\_31 는 모든 분위 수에서  
단일 값을 가진다고 판단됨

Verification\_ Q-Q graph for some Features

insight

R\_26



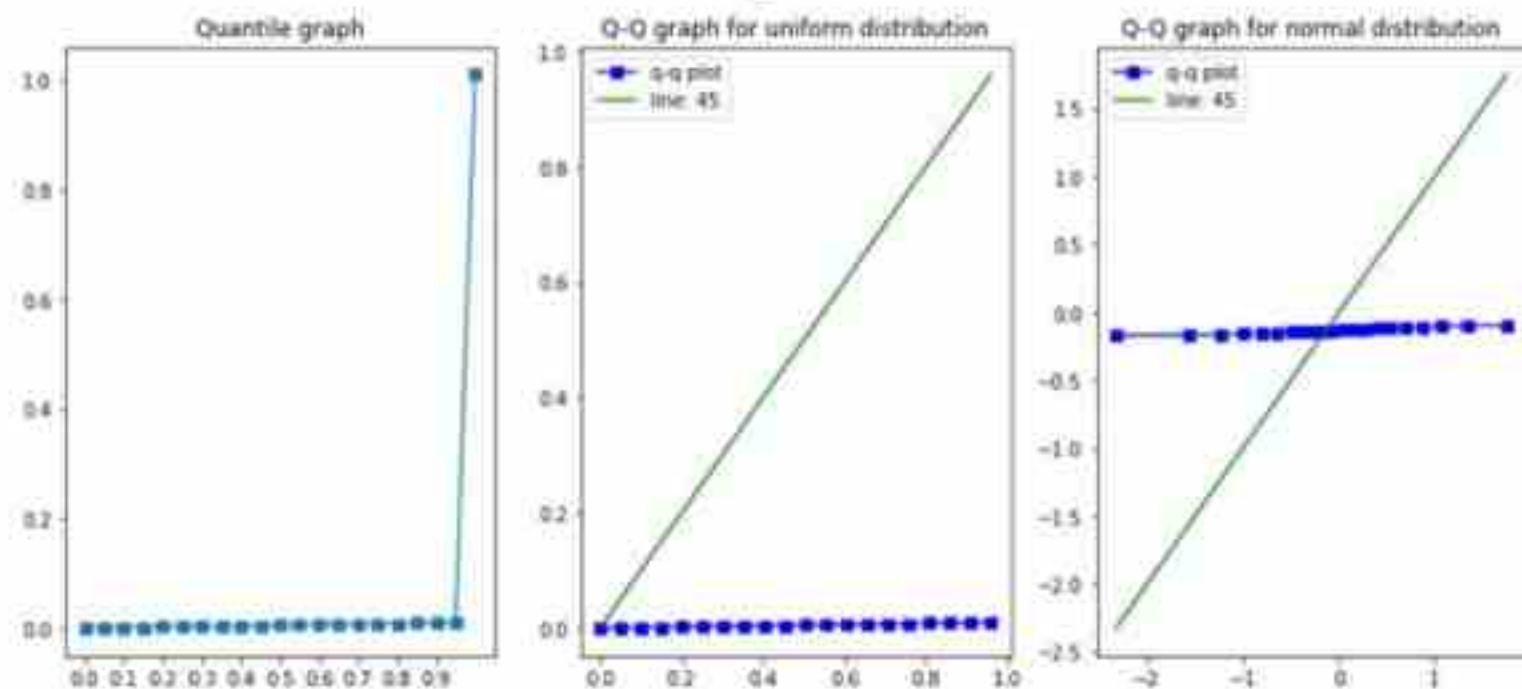
R\_26 는 확대 시 6개의  
discrete value가 있는 것  
으로 판단됨



Verification\_ Q-Q graph for some Features

insight

D\_94



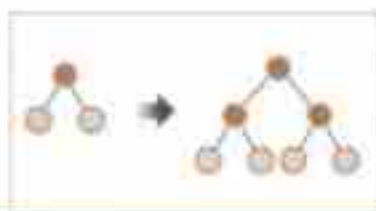
D\_94, S\_26 기울기 거의 0  
→ 이상치 값이 매우 크다고 판단됨

# 3



**Insight about Machine Learning Model**

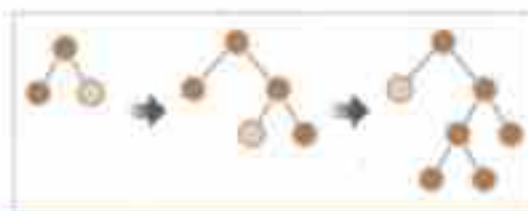
## Part 3 >> Generate Machine Learning Model



### 1 Random Forest

ensemble ⊃ bagging ⊃ RF

- subsample에 대해 다수의 결정 트리 classifier를 최적화
- 여러 추론 결과 평균 → 예측 정확도를 개선 → 과적합방지
- hyperparameter xgbm보다 적음



### 2 XGBoost

Ensemble ⊃ boosting ⊃ GBM ⊃ XGBoost, LightGBM

- 트리기반 앙상블 균형트리분할방식
- 병렬학습이 가능 → GBM의 단점인 오랜 수행시간을 극복
- 예측 성능 우수
- 가지치기 기능 존재 조기종단 기능 존재

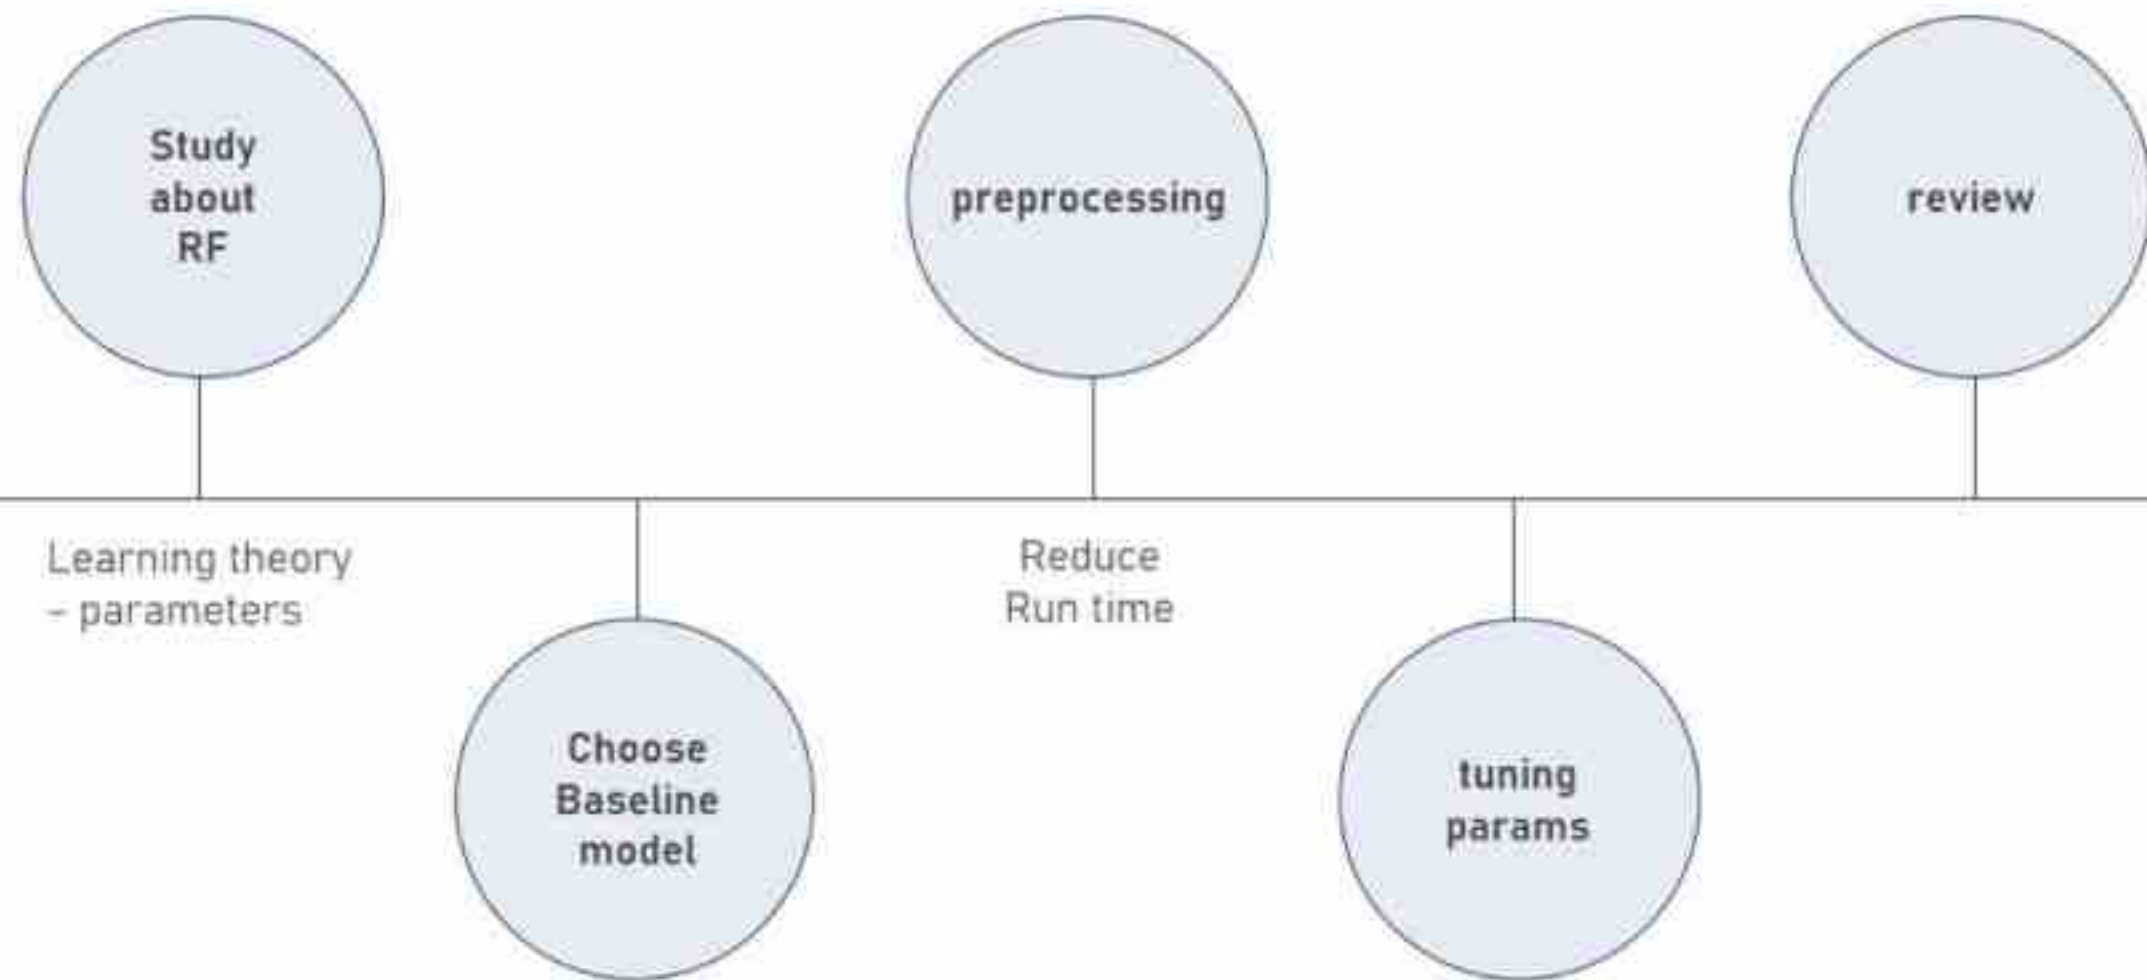
### 3 LightGBM

- 트리기반 앙상블 리프중심트리분할방식
- XGBoost보다 더 빠르고 메모리 사용량 적음
- 예측 성능 우수
- 10000건 이하의 dataset 처리시 과적합 가능성 ↑

### 4 CatBoost

- Categorical Boosting 범주형 feature 처리중점
- Gradient Boosting 기반 XGB, LGBM보다 우수
- Feature 자동 타깃인코딩
- 망각 결정 트리 사용 (oblivious decision tree)


## Part 3 >> Generate Machine Learning Model



## Part 3 >> Generate Machine Learning Model\_ Random Forest

### Missing Value Processing\_ Remove Columns

```
In [12]: #lets remove columns if there are >90% of missing values.  
#Given that there are many columns with large number of missing values, it is impractical to go through every single one of them to determine whether it is useful.  
#Furthermore, we do not have information on the feature (e.g. actual name of the feature) except the type of variable  
#Brute force is thus a practical option to weed out columns with too many missing values.  
#Since about 33.1% of the data are defaults(66.9% non-defaults), it is safe to say that columns with  
>90% missing data are not useful.  
train=train_data.dropna(axis=1, thresh=int(0.90*len(train_data)))  
  
#Checking the shape of new train data  
train.shape  
## We are now left with 152 columns  
  
Out[12]:  
(5531451, 152)
```



191 columns → 152 columns

## Part 3 >> Generate Machine Learning Model\_ Random Forest

### Missing Value Processing\_ Fill Missing Values, Drop S\_2 groupby

```
In [10]:  
# There are multiple transactions. Lets take only the latest transaction from each customer.  
# latest transaction may have missing values, we will perform forward fill for these missing values.  
# We perform forward fill as the last known value is likely to be brought forward to the next transaction.  
# We then do a backfill if the first row happens to be NA.  
train=train.set_index(['customer_ID'])  
train=train.ffill().bfill()  
train=train.reset_index()  
train=train.groupby('customer_ID').tail(1)  
train=train.set_index(['customer_ID'])  
  
#Drop S2 column since it is no longer relevant  
train=train.drop(['S2'],axis=1,inplace=True)  
#Check the number of rows  
train.shape  
# We now have 458913 rows, which corresponds to the number of unique customers.  
  
Out[11]:  
(458913, 150)
```

Rows : 5531451 → 458913

Columns : 152 → 150



one-hot Encoding for D\_63 and D\_64

- 문자열인 범주형 변수를 원핫인코딩

```
train_D63 = pd.get_dummies(train[['D_63']])  
train = pd.concat([train, train_D63], axis=1)  
train = train.drop(['D_63'], axis=1)  
  
train_D64 = pd.get_dummies(train[['D_64']])  
train = pd.concat([train, train_D64], axis=1)  
train = train.drop(['D_64'], axis=1)
```

### Remove Highly Correlated Features

- 상관관계 높은 feature 제거
- absolute correlation >= 90% 이상인 열 drop

```
train_without_target=train.drop(['target'],axis=1)
cor_matrix = train_without_target.corr().abs()
upper_tri = cor_matrix.where(np.triu(np.ones(cor_matrix.shape),k=1).astype(np.bool))
#Drop out columns with absolute correlation of more than 90%
to_drop = [column for column in upper_tri.columns if any(upper_tri[column] > 0.90)]
train_drop_highcorr=train.drop(to_drop,axis=1)
train_drop_highcorr.shape
#We are now left with 145 columns, which is still significant
```

(458913, 145)

Columns : 150 → 145

### Remove Columns with Low Variance = 0.1

```
from sklearn.feature_selection import VarianceThreshold
from itertools import compress
def fs_variance(df, threshold=float(0.1)):
    features = list(df.columns)

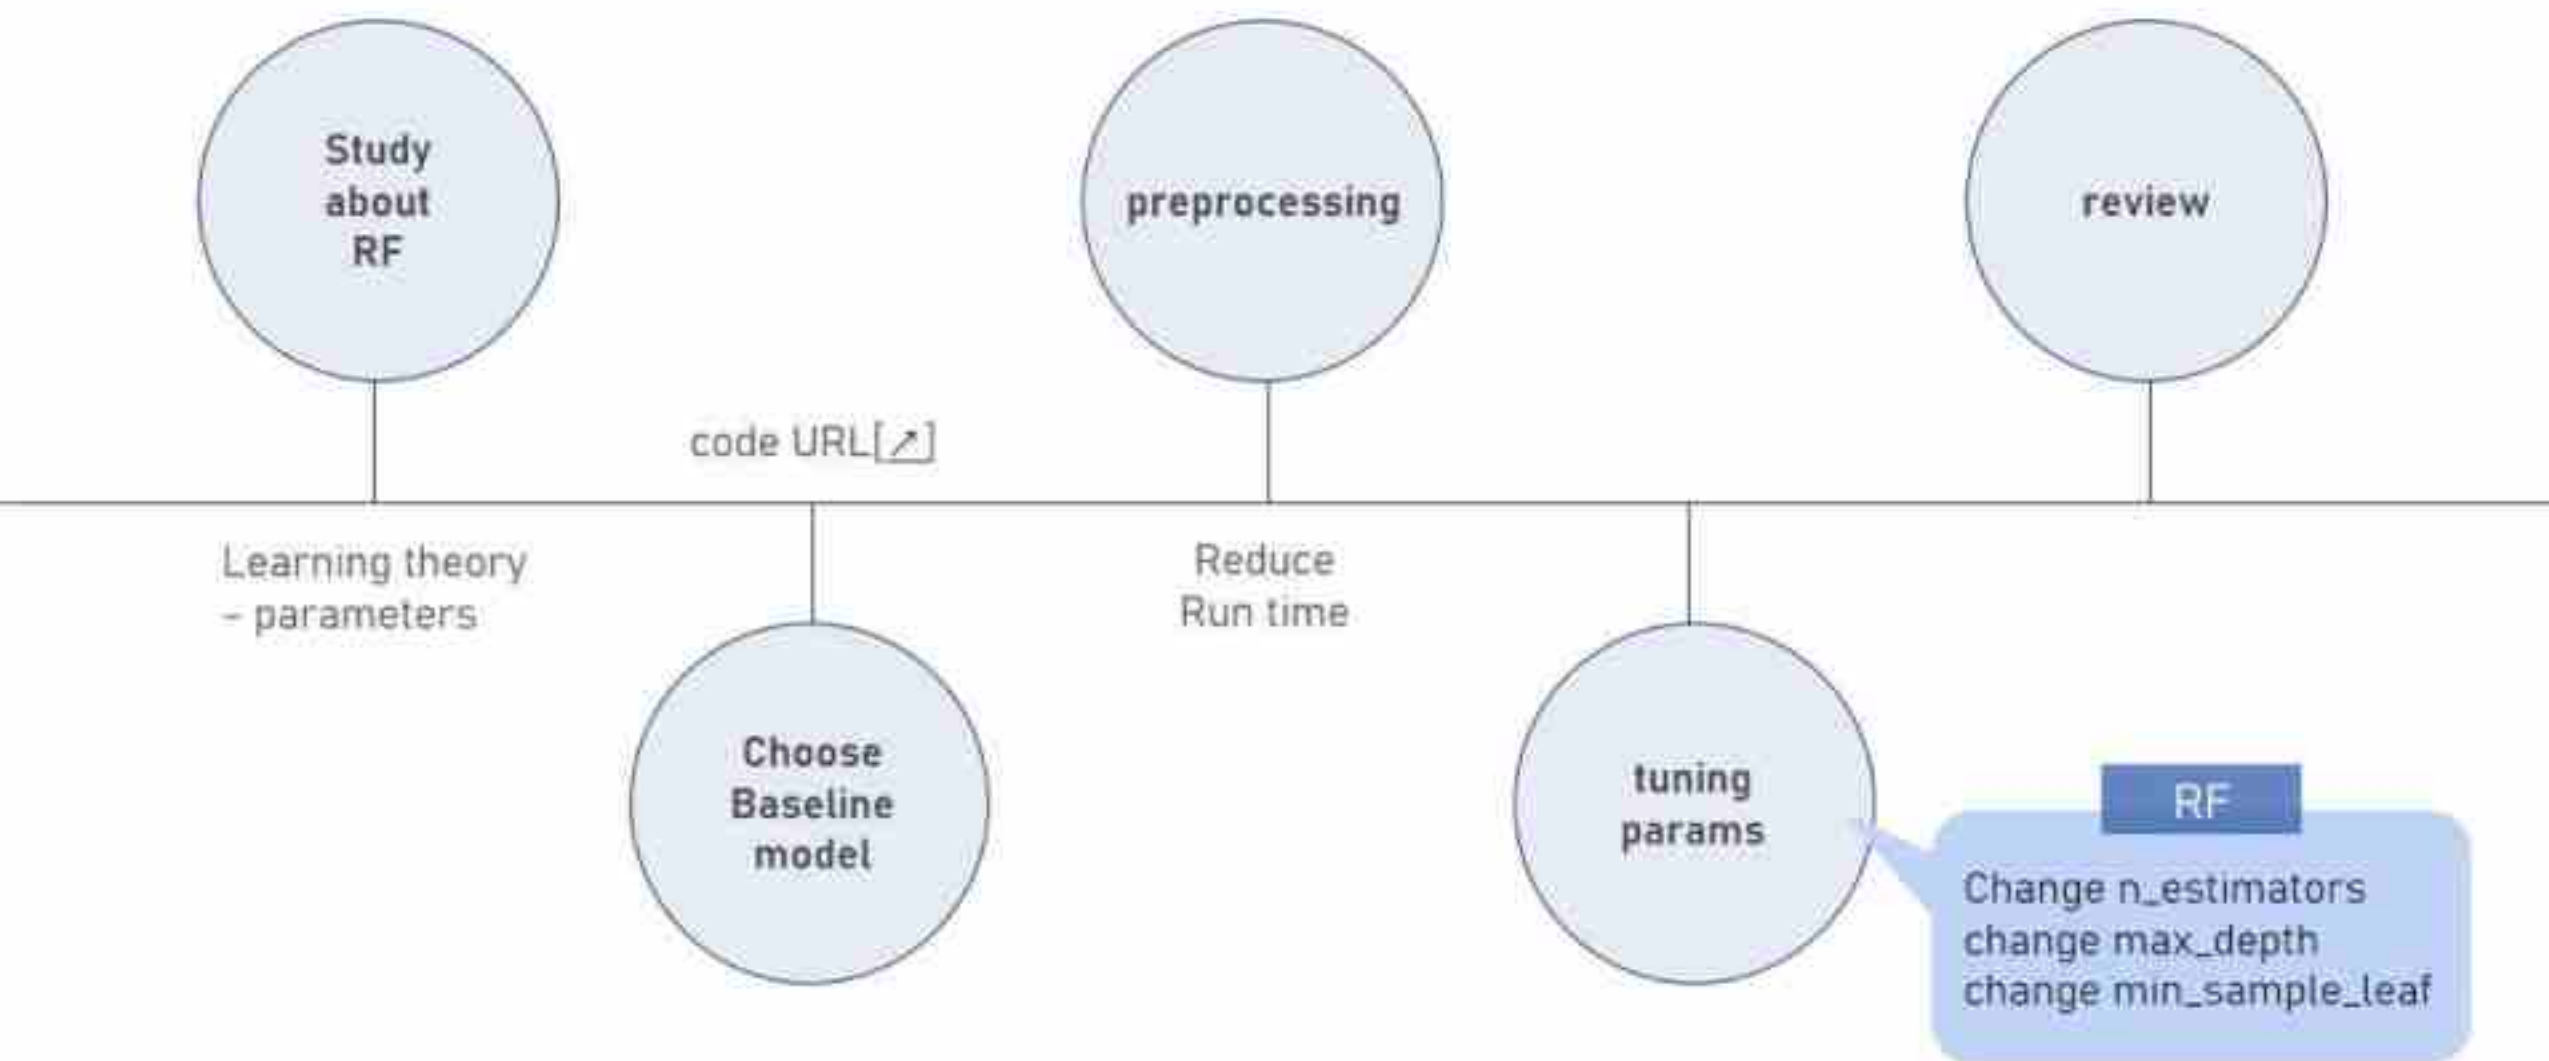
    # Initialize and fit the method
    vt = VarianceThreshold(threshold = threshold)
    _ = vt.fit(df)

    # Get which column names which pass the threshold
    feat_select = list(compress(features, vt.get_support()))

    return feat_select
columns_to_keep=fs_variance(train.drop_highcorr)
# We are left with 54 columns (excluding target), which passed the threshold.
train_final=train[columns_to_keep]
len(columns_to_keep)
```

54 columns left (except target)

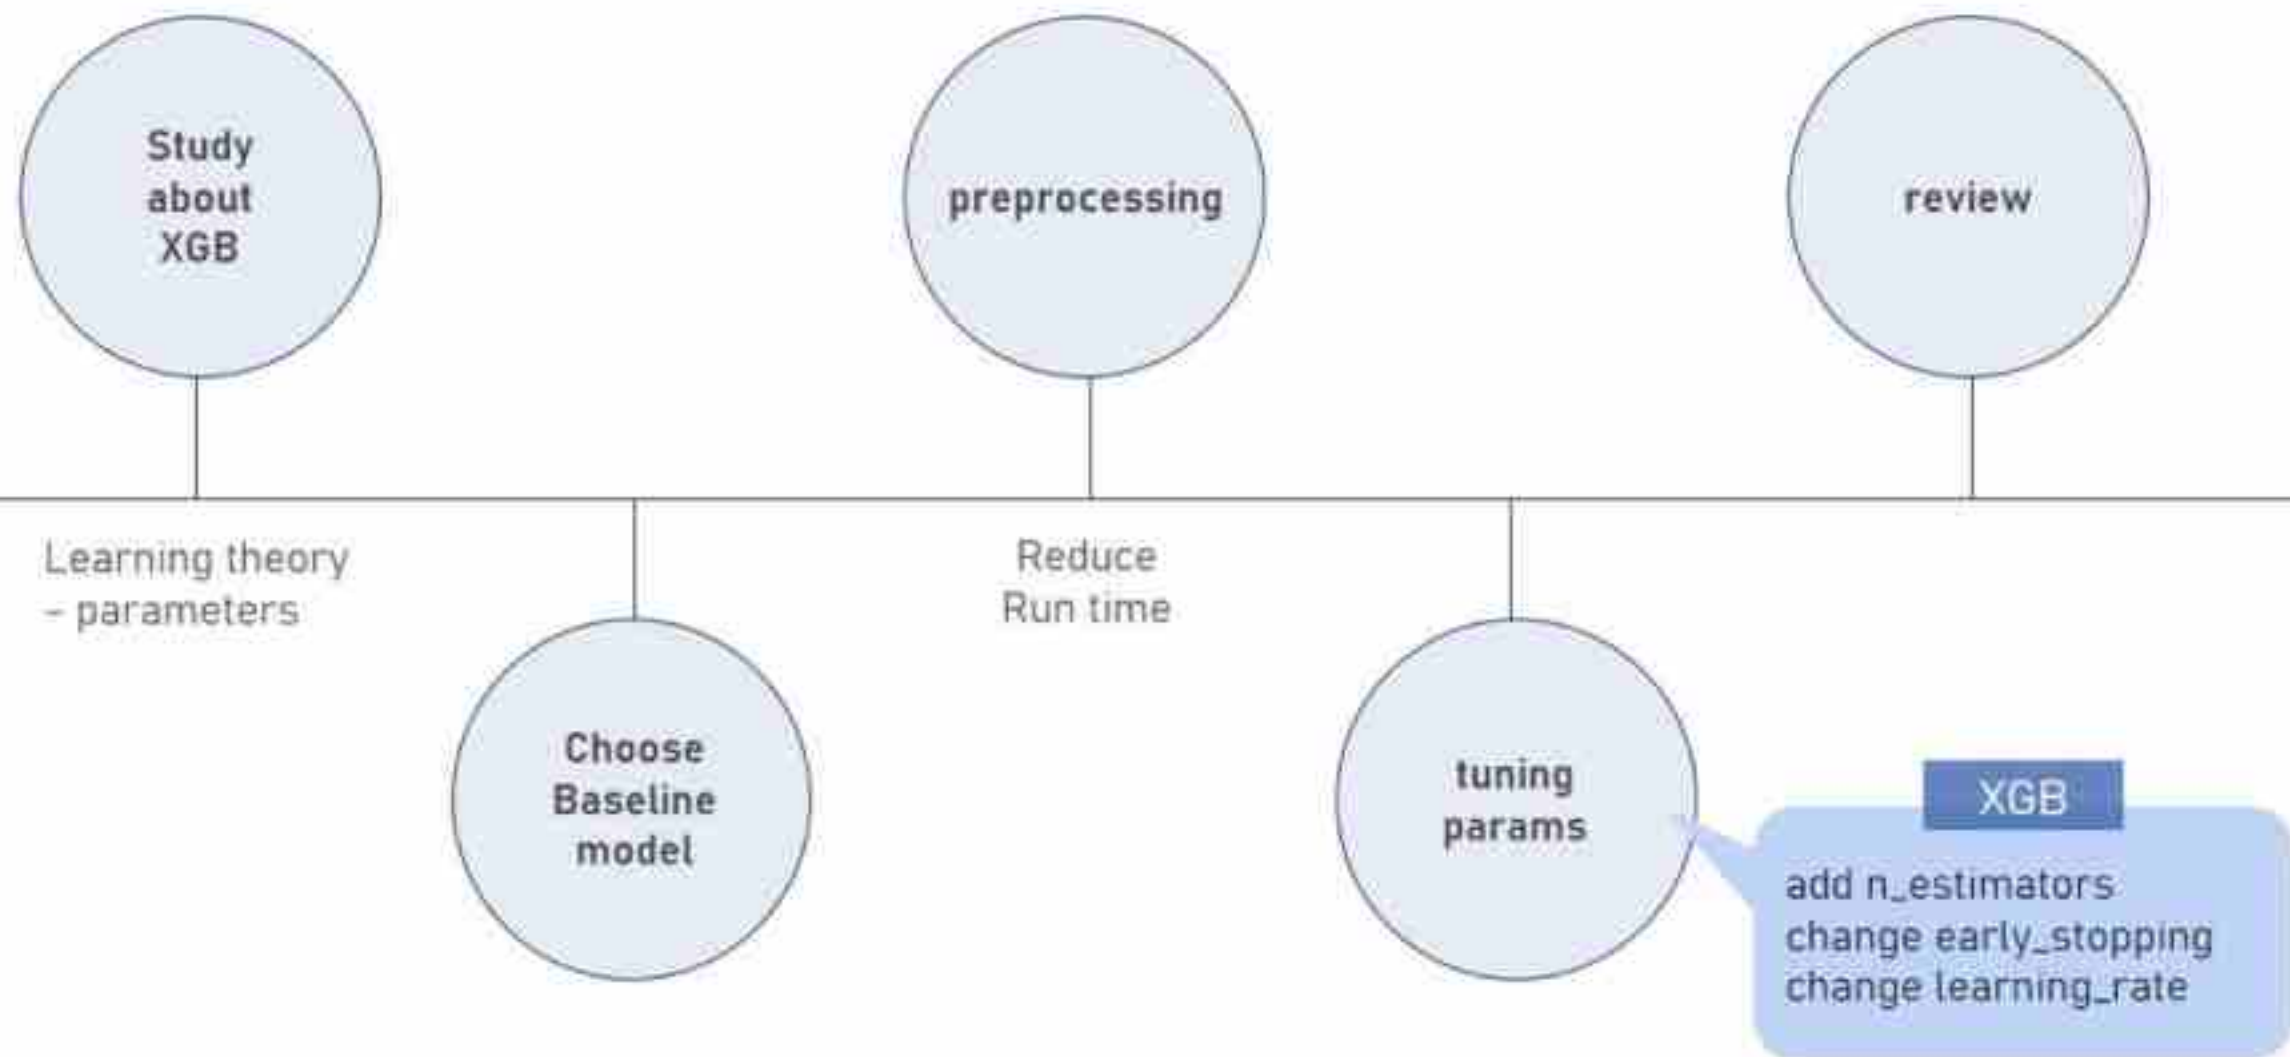
## Part 3 >> Generate Machine Learning Model\_ Random Forest



## Part 3 >> Generate Machine Learning Model\_ Random Forest

No.	내용요약	소요시간	parameters	scores
1	Base RandomForest model	1058.8s - GPU	n_estimators=400, max_features='sqrt', bootstrap=True(default), max_depth=30, min_samples_leaf=1(default), min_samples_split=5, n_jobs=-1	score: 0.729
2		477.1s - CPU		
3	change n_estimators (v1)	290.8s - CPU	n_estimators → 200	score: 0.727 / ↓
4		437.6s - CPU	n_estimators → 500	score: 0.729 / -
5	change max_depth (v2)	547.8s - CPU	max_depth → 100	score: 0.728 / ↓
6		538.3s - CPU	max_depth → 40	score: 0.729 / -
7		535.3s - CPU	max_depth → 50	score: 0.729 / -
8		244.6s - CPU	max_depth → 10	score: 0.720 / ↓
9		303.5s - CPU	max_depth → 20	score: 0.742 / ↑
10	add random_state (v3)	383.9s - CPU	random_state = 77	score: 0.730 / ↑
11	change min_samples_leaf (v4)	349.5s - CPU	min_samples_leaf → 2	score: 0.744 / ↑

## Part 3 >> Generate Machine Learning Model\_ XGBoost





## Part 3 >> Generate Machine Learning Model\_ XGBoost

No.	내용요약	소요시간	parameters	scores
1	basic XGB model	190.8s(GPU)	(n_estimators : default 100, early_stopping_rounds=100; learning_rate=0.05)	변동없음
2	add n_estimators	186.6s(GPU)	n_estimators=110	
3		176.8s(GPU)	n_estimators=120	
4		190.4s(GPU)	n_estimators=130	
5		191.8s(GPU)	n_estimators=140	
6		196.2s(GPU)	n_estimators=150	
7		183.4s - GPU	n_estimators=200	
8	change early_stopping option	192.0s - GPU	n_estimators=200, early_stopping_rounds=150	
9	change n_estimators	201.1s - GPU	n_estimators=300, early_stopping_rounds=150	
10	change learning_rate option	179.0s - GPU	n_estimators=200, learning_rate=0.1, early_stopping_rounds=100	
11		246.3s - GPU	n_estimators=200, learning_rate=0.2, early_stopping_rounds=100	

Find out the Best Parameters\_Define Parameter Range

In [13]:

```
#define parameter range  
learning_rate=np.linspace(0.01,0.1,10)  
max_depth=np.arange(2, 18, 2)  
colsample_bylevel=np.arange(0.3, 0.8, 0.1)  
iterations=np.arange(50, 1000, 50)  
l2_leaf_reg=np.arange(0,10)  
bagging_temperature=np.arange(0,100,10)  
n_estimators=np.arange(50,500,50)
```

### Define Parameter Space, Fit Condition, and Ross Function

```
xgb_cat_params = {
    'learning_rate': hp.choice('learning_rate', learning_rate),
    'max_depth': hp.choice('max_depth', max_depth),
    'colsample_bytree': hp.choice('colsample_bytree', colsample_bylevel),
    'n_estimators': hp.choice('n_estimators', n_estimators),
    'loss_function': 'logloss',
    'nan_mode': 'Min',
    'task_type': 'GPU'
}
xgb_fit_params = {
    'eval_metric': 'logloss',
    'early_stopping_rounds': 10,
    'verbose': False
}
xgb_para = dict()
xgb_para['cls_params'] = xgb_cat_params
xgb_para['fit_params'] = xgb_fit_params
xgb_para['loss_func'] = lambda y, pred: np.sqrt(mean_squared_error(y, pred))
```

### Define Parameter Space, Fit Condition, and Ross Function

```
#calling the XGBoost function by passing XGB parameter space
xgb_opt = obj.process(fn_name='xgb_cls', space=xgb_para, trials=Trials(), algo=tpe.suggest,
max_evals=2)

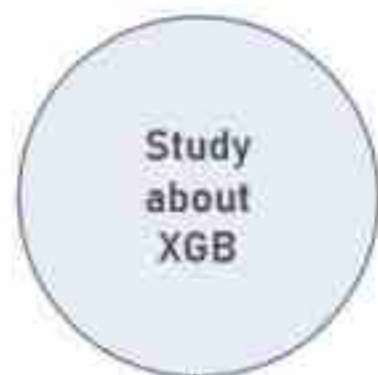
In [24]:
#Save best parameters in a dictionary
best_param_xgb={}
best_param_xgb['learning_rate']=learning_rate[xgb_opt['learning_rate']]
best_param_xgb['colsample_bytree']=colsample_bytree[xgb_opt['colsample_bytree']]
best_param_xgb['max_depth']=max_depth[xgb_opt['max_depth']]
best_param_xgb['n_estimators']=n_estimators[xgb_opt['n_estimators']]

In [25]:
best_param_xgb

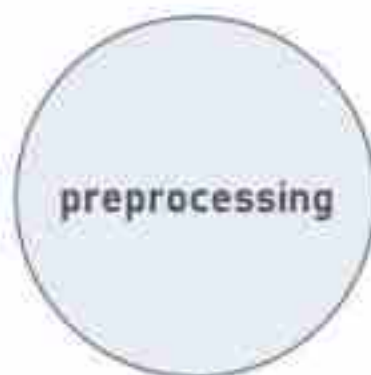
Out[25]:
{'learning_rate': 0.05000000000000001,
 'colsample_bytree': 0.6000000000000001,
 'max_depth': 8,
 'n_estimators': 400}
```

## Part 3 >> Generate Machine Learning Model\_ XGBoost

F1-score of XGBoost : 0.81



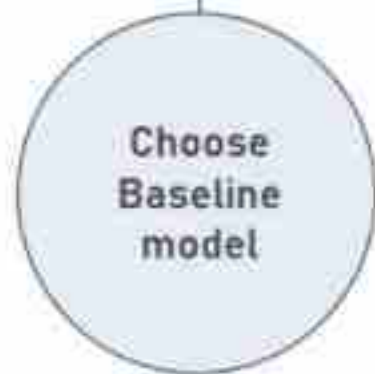
Learning theory  
- parameters



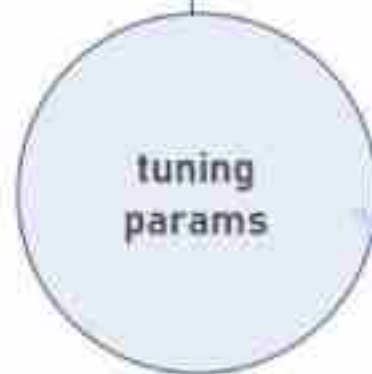
Reduce  
Run time



tuning  
params



Choose  
Baseline  
model



**XGB**

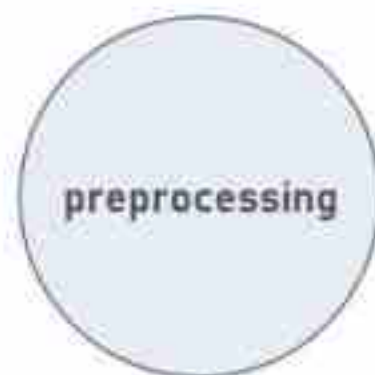
n\_estimators = 400  
colsample\_bytree = 0.6  
max\_depth = 8  
learning\_rate = 0.05

## Part 3 >> Generate Machine Learning Model\_ LightGBM

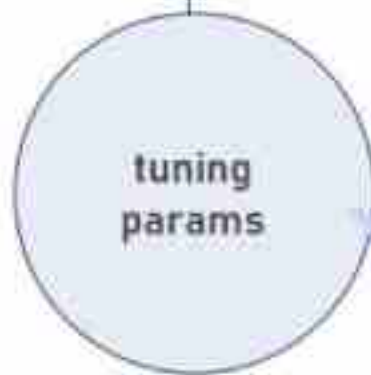
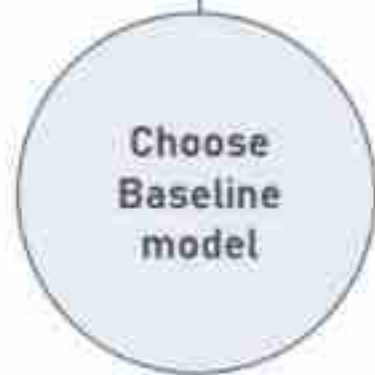
F1-score of LGBM : 0.81



Learning theory  
- parameters



Reduce  
Run time



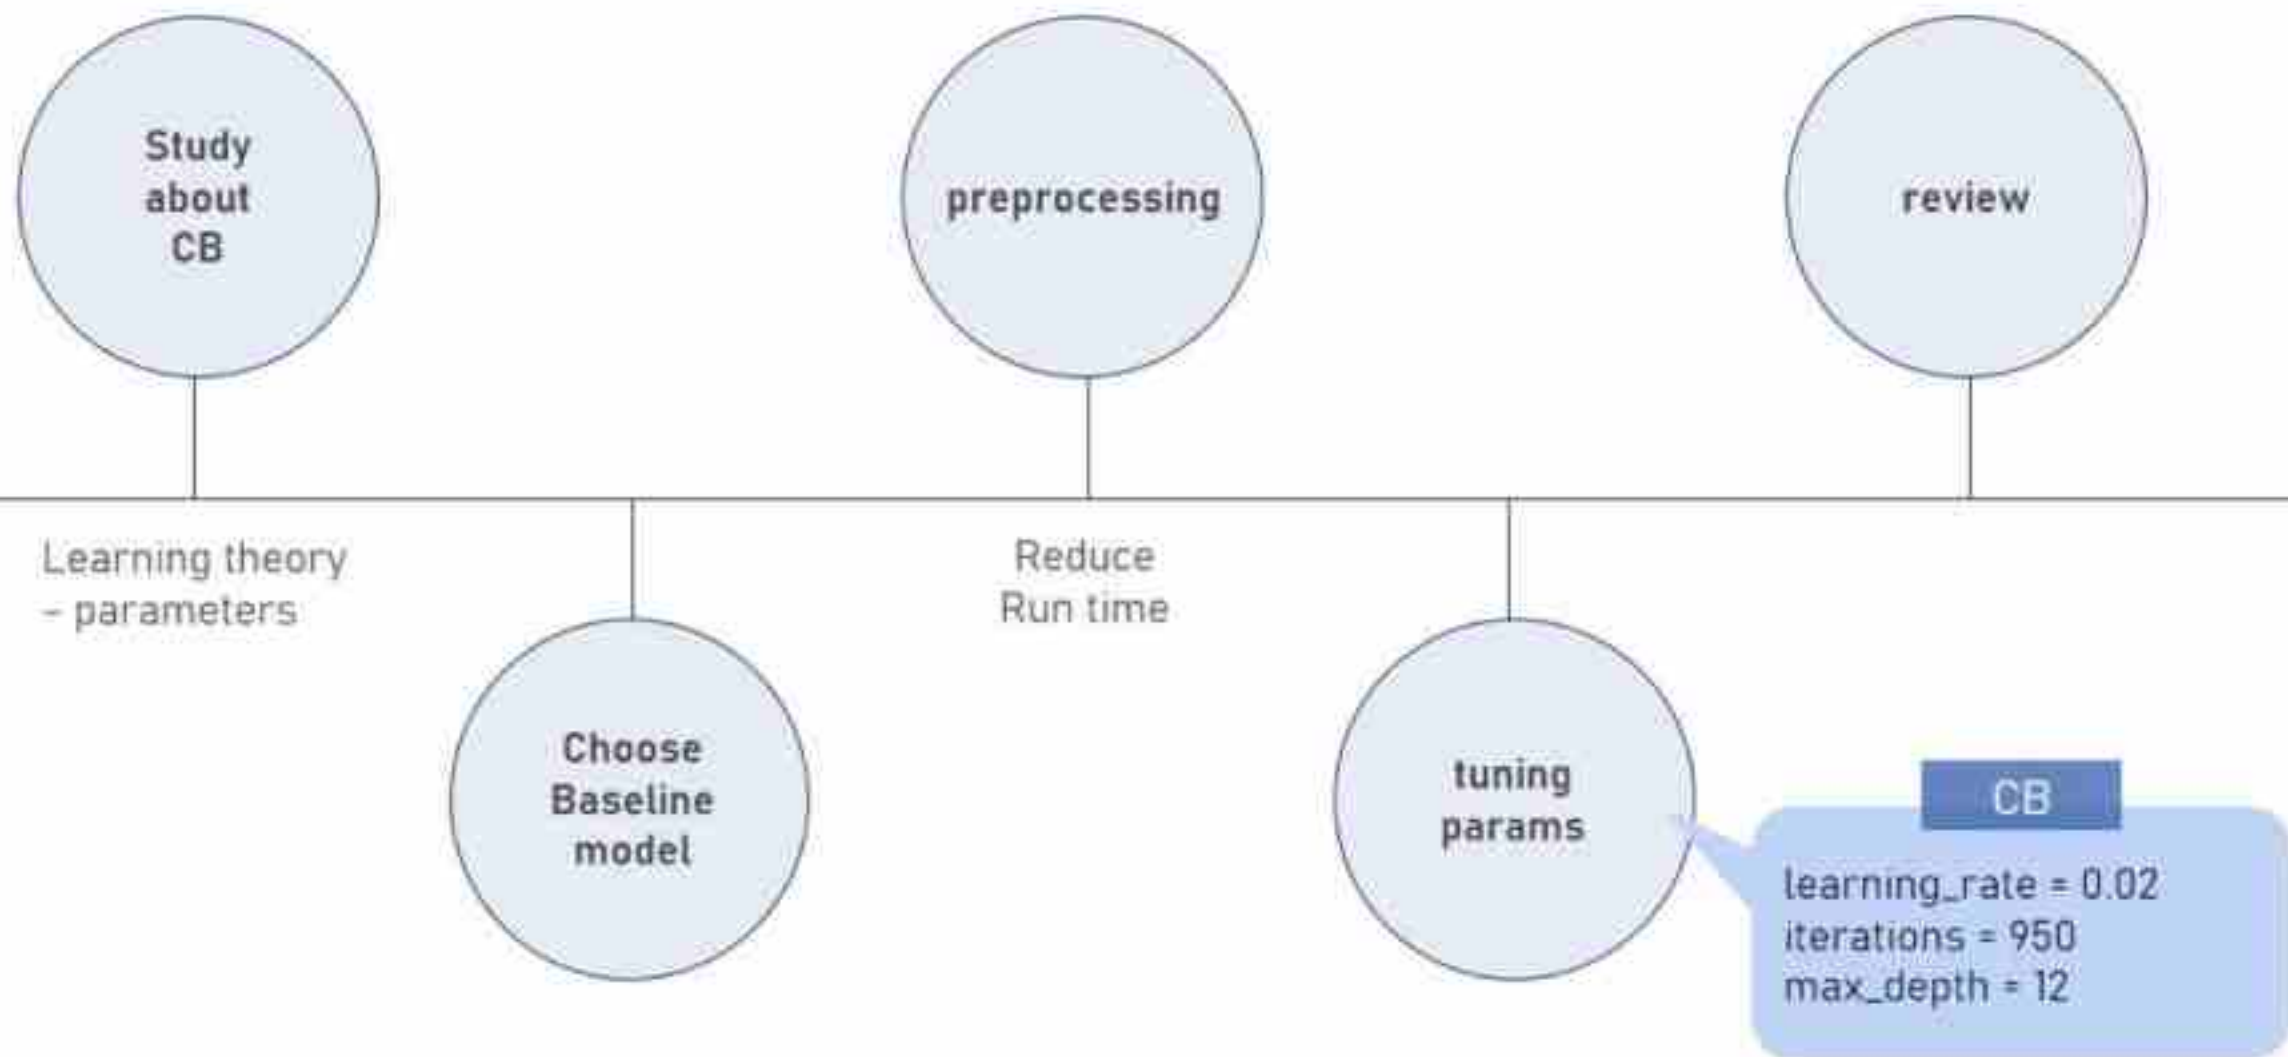
**LGBM**

n\_estimators = 400  
colsample\_bytree = 0.6  
max\_depth = 14  
learning\_rate = 0.04



## Part 3 >> Generate Machine Learning Model\_ CatBoost

F1-score of Catboost: 0.80



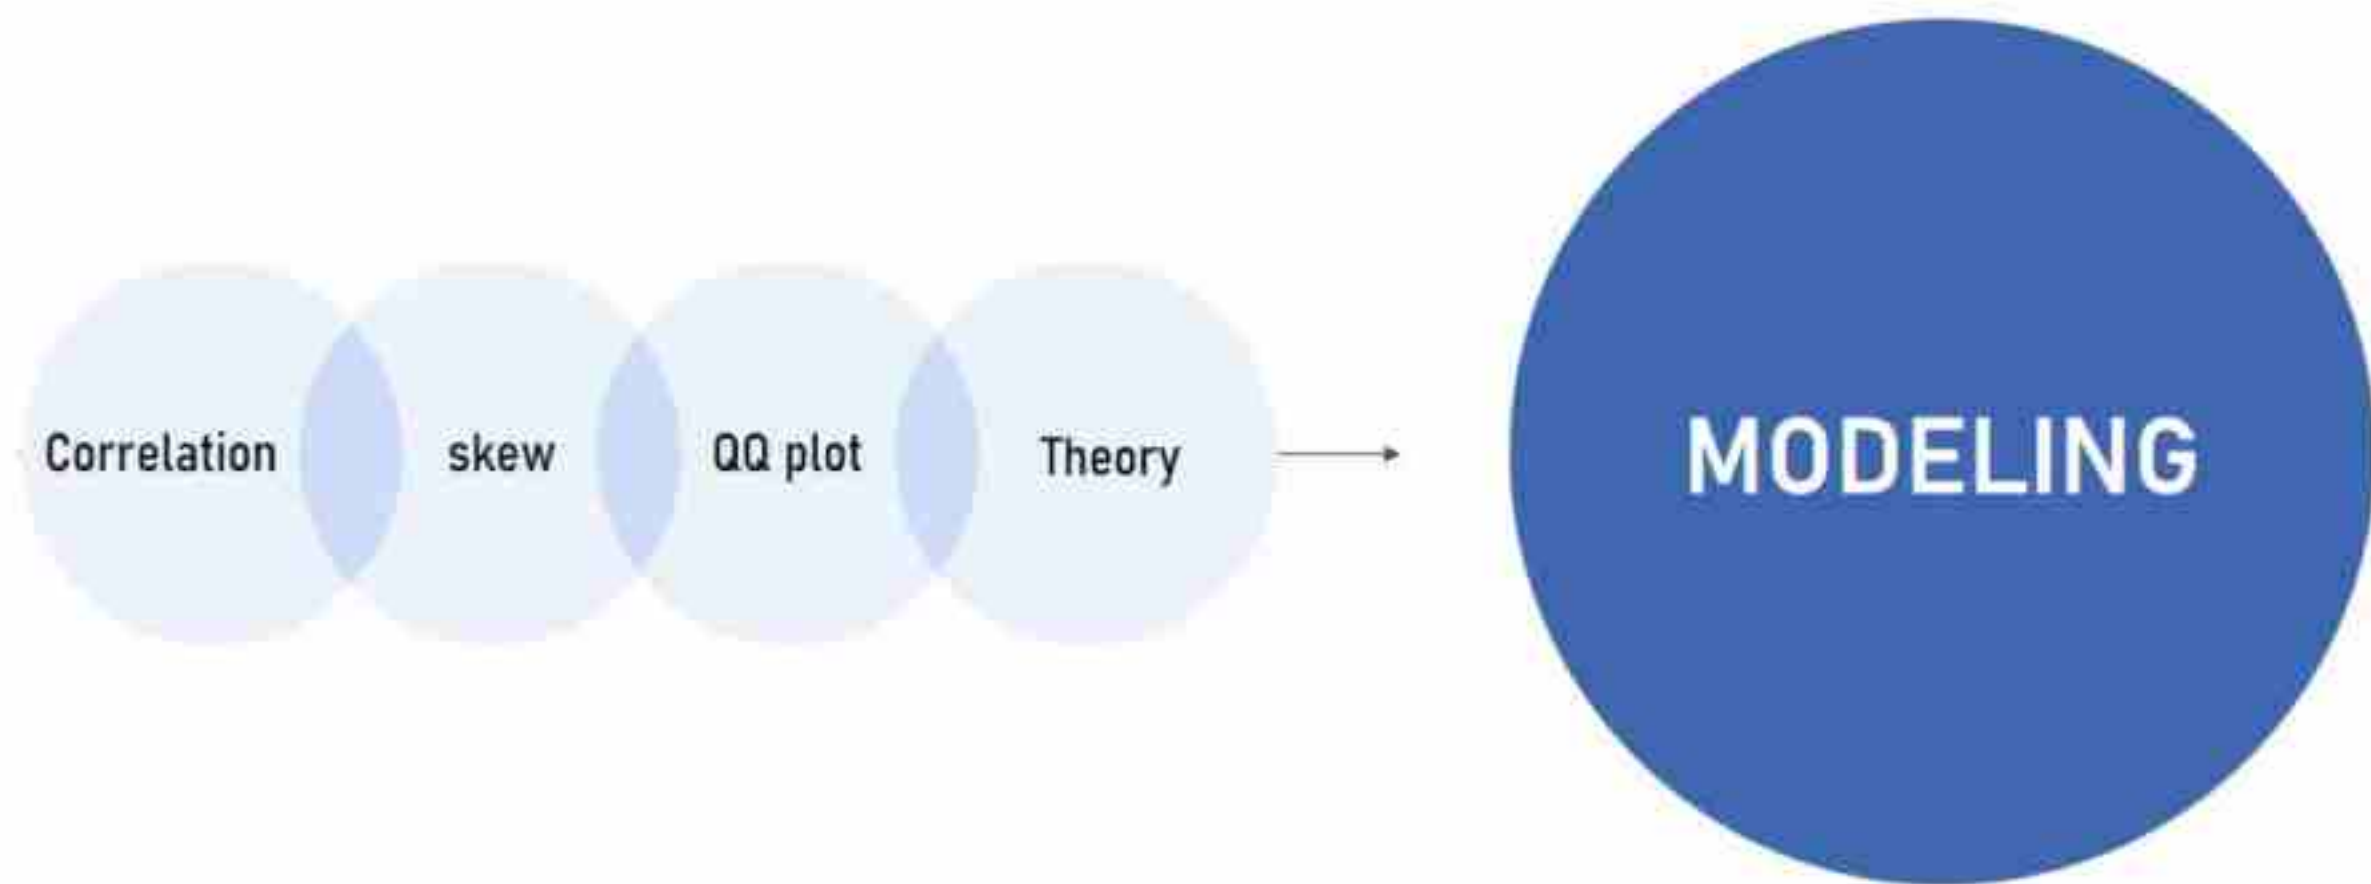


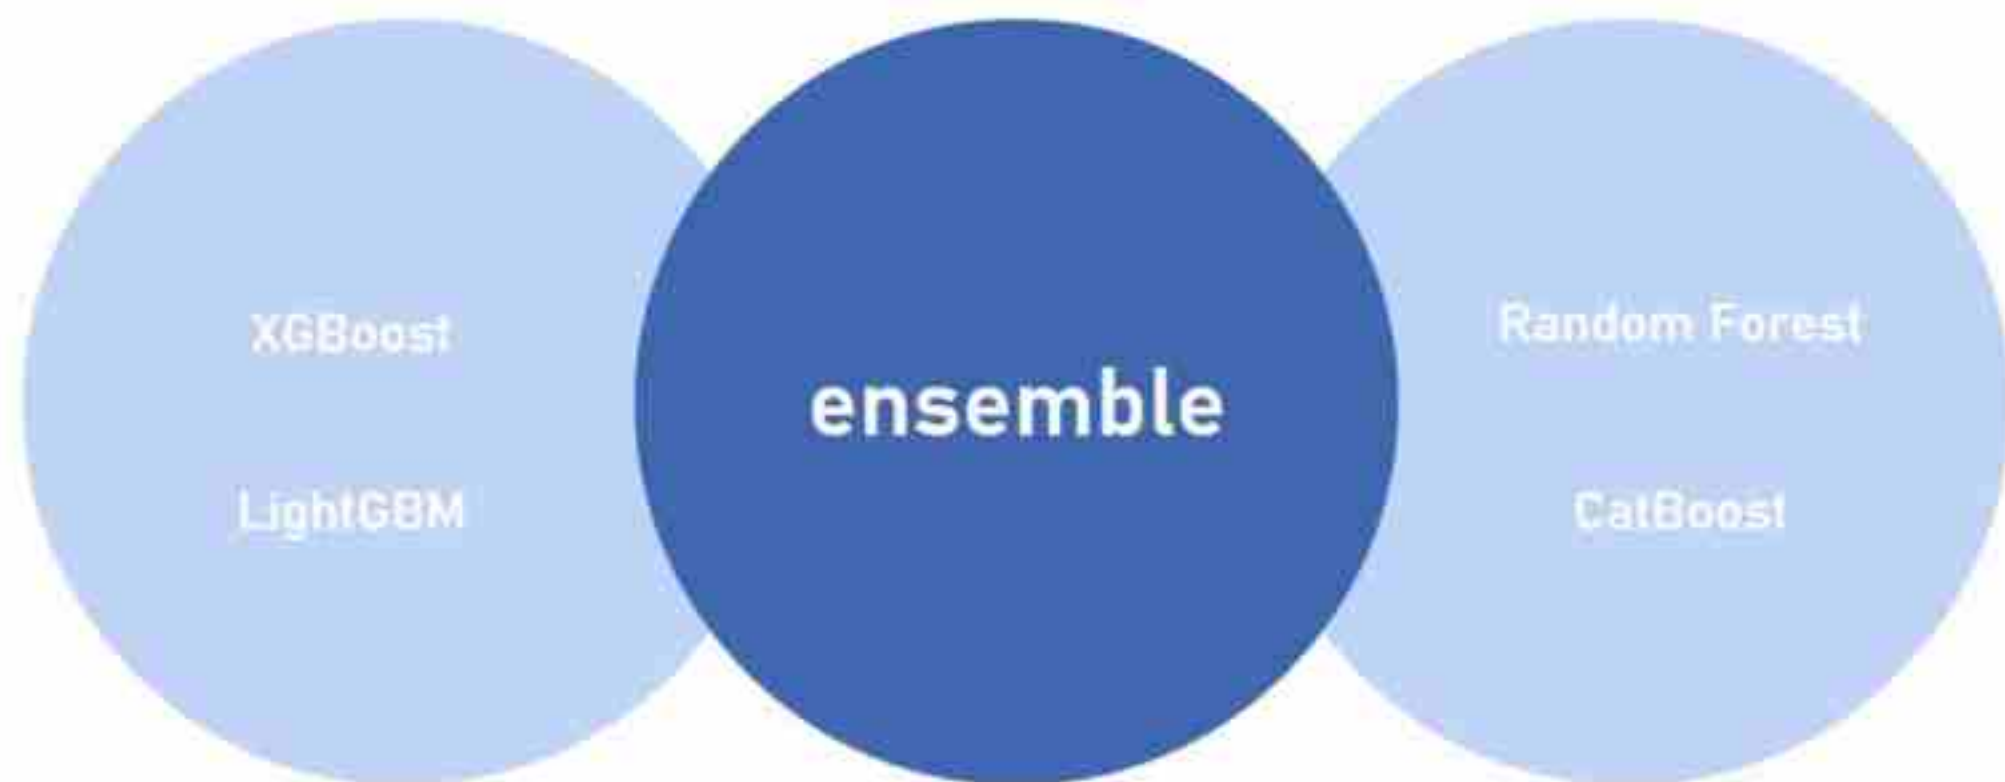
# 4



## Conclusion

## Part 4 >> Conclusion





A dimly lit, modern kitchen interior. In the center, there is a wooden kitchen island with a light-colored countertop. Two white, modern-style bar stools with chrome bases are positioned in front of the island. The background shows a dark wall and a window with a view of the outdoors. The overall atmosphere is quiet and contemporary.

**Q&A**